



> Home > > Microsoft FastTrack > > FastTrack for Azure
> > Create an Azure OpenAI, LangChain, ChromaDB, and Chainlit Chat App
in Container Apps using Terraform

[Back to Blog](#)[< Newer Article](#)[Older Article >](#)

Create an Azure OpenAI, LangChain, ChromaDB, and Chainlit Chat App in Container Apps using Terraform

...

By  Paolo Salvatori

Published Jul 27 2023 06:47 AM

 17.8K Views



This article shows how to quickly build chat applications using [Python](#) and leveraging powerful technologies such as [OpenAI ChatGPT models](#), [Embedding models](#), [LangChain](#) framework, [ChromaDB](#) vector database, and [Chainlit](#), an open-source Python package that is specifically designed to create user interfaces (UIs) for AI applications. These applications are hosted on [Azure Container Apps](#), a fully managed environment that enables you to run microservices and containerized applications on a serverless platform.

- **Simple Chat:** This simple chat application utilizes OpenAI's language models to generate real-time completion responses.
- **Documents QA Chat:** This chat application goes beyond simple conversations. Users can upload up to 10 .pdf and .docx documents, which are then processed to create vector embeddings. These embeddings are stored in [ChromaDB](#) for efficient retrieval. Users can pose questions about the uploaded documents and view the [Chain of Thought](#), enabling easy exploration of the reasoning process. The completion message contains links to the text chunks in the documents that were used as a source for the response.

Both applications use a [user-defined managed identity](#) to authenticate and authorize against [Azure OpenAI Service \(AOAI\)](#), and [Azure Container Registry \(ACR\)](#), and use [Azure Private Endpoints](#) to connect privately and securely to these services. The chat UIs are built using [Chainlit](#), an open-source Python package designed

explicitly for creating AI applications. Chainlit seamlessly integrates with [LangChain](#), [LlamaIndex](#), and [LangFlow](#), making it a powerful tool for easily developing ChatGPT-like applications.

By following our example, you can quickly create sophisticated chat applications that utilize cutting-edge technologies, empowering users with intelligent conversational capabilities.

You can find the code and Visio diagrams in the companion [GitHub](#) repository. Also, check the following articles:

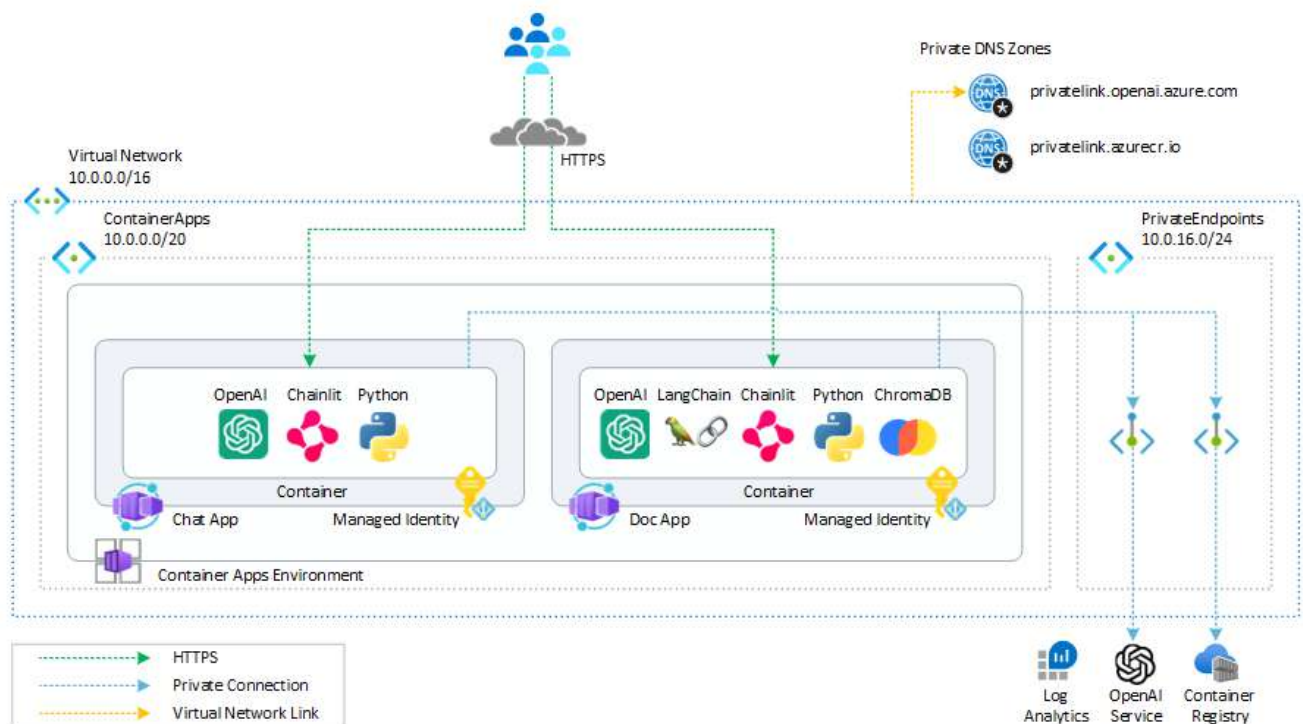
- [Deploy and run an Azure OpenAI ChatGPT application on AKS via Bicep](#)
- [Deploy and run an Azure OpenAI ChatGPT application on AKS via Terraform](#)

Prerequisites

- An active [Azure subscription](#). If you don't have one, create a [free Azure account](#) before you begin.
- [Visual Studio Code](#) installed on one of the [supported platforms](#) along with the [HashiCorp Terraform](#).
- Azure CLI version 2.49.0 or later installed. To install or upgrade, see [Install Azure CLI](#).
- aks-preview Azure CLI extension of version 0.5.140 or later installed
- [Terraform v1.5.2 or later](#).

Architecture

The following diagram shows the architecture and network topology of the sample:



This sample provides two sets of Terraform modules to deploy the infrastructure and the chat applications.

Infrastructure Terraform Modules

You can use the Terraform modules in the `terraform/infra` folder to deploy the infrastructure used by the sample, including the [Azure Container Apps Environment](#), [Azure OpenAI Service \(AOAI\)](#), and [Azure Container Registry \(ACR\)](#), but not the [Azure Container Apps \(ACA\)](#). The Terraform modules in the `terraform/infra` folder

[Skip to later navigation](#)

deploy the following resources:

- azurerm_virtual_network: an Azure Virtual Network with two subnets:
 - ContainerApps : this subnet hosts the Azure Container Apps Environment.
 - PrivateEndpoints : this subnet contains the Azure Private Endpoints to the Azure OpenAI Service (AOAI) and Azure Container Registry (ACR) resources.
- azurerm_container_app_environment: the Azure Container Apps Environment hosting the Azure Container Apps.
- azurerm_cognitive_account: an Azure OpenAI Service (AOAI) with a GPT-3.5 model used by the chatbot applications. Azure OpenAI Service gives customers advanced language AI with OpenAI GPT-4, GPT-3, Codex, and DALL-E models with Azure's security and enterprise promise. Azure OpenAI co-develops the APIs with OpenAI, ensuring compatibility and a smooth transition from one to the other. The Terraform modules create the following models:
 - GPT-35: a gpt-35-turbo-16k model is used to generate human-like and engaging conversational responses.
 - Embeddings model: the text-embedding-ada-002 model is to transform input documents into meaningful and compact numerical representations called embeddings. Embeddings capture the semantic or contextual information of the input data in a lower-dimensional space, making it easier for machine learning algorithms to process and analyze the data effectively. Embeddings can be stored in a vector database, such as ChromaDB or Facebook AI Similarity Search, explicitly designed for efficient storage, indexing, and retrieval of vector embeddings.
- azurerm_user_assigned_identity: a user-defined managed identity used by the chatbot applications to acquire a security token to call the Chat Completion API of the ChatGPT model provided by the Azure OpenAI Service and to call the Embedding model.
- azurerm_container_registry: an Azure Container Registry (ACR) to build, store, and manage container images and artifacts in a private registry for all container deployments. In this sample, the registry stores the container images of the two chat applications.
- azurerm_private_endpoint: an Azure Private Endpoint is created for each of the following resources:
 - Azure OpenAI Service (AOAI).
 - Azure Container Registry (ACR).
- azurerm_private_dns_zone: an Azure Private DNS Zone is created for each of the following resources:
 - Azure OpenAI Service (AOAI).
 - Azure Container Registry (ACR).
- azurerm_log_analytics_workspace: a centralized Azure Log Analytics workspace is used to collect the diagnostics logs and metrics from all the Azure resources:
 - Azure OpenAI Service (AOAI).
 - Azure Container Registry (ACR).
 - Azure Container Apps (ACA).

Application Terraform Modules

You can use these Terraform modules in the `terraform/apps` To deploy the Azure Container Apps (ACA) using the Docker container images stored in the Azure Container Registry you deployed in the previous step.

- azurerm_container_app: this sample deploys the following applications:
 - **chatapp**: this simple chat application utilizes OpenAI's language models to generate real-time completion responses.

~~Skip to final configuration~~

- **docapp:** This chat application goes beyond conversations. Users can upload up to 10 .pdf and .docx documents, which are then processed to create vector embeddings. These embeddings are stored in ChromaDB for efficient retrieval. Users can pose questions about the uploaded documents and view the [Chain of Thought](#), enabling easy exploration of the reasoning process. The completion message contains links to the text chunks in the files that were used as a source for the response.

Azure Container Apps

[Azure Container Apps \(ACA\)](#) is a serverless compute service provided by Microsoft Azure that allows developers to easily deploy and manage containerized applications without the need to manage the underlying infrastructure. It provides a simplified and scalable solution for running applications in containers, leveraging the power and flexibility of the Azure ecosystem.

With Azure Container Apps, developers can package their applications into containers using popular containerization technologies such as [Docker](#). These containers encapsulate the application and its dependencies, ensuring consistent execution across different environments.

Powered by [Kubernetes](#) and open-source technologies like [Dapr](#), [KEDA](#), and [envoy](#), the service abstracts away the complexities of managing the infrastructure, including provisioning, scaling, and monitoring, allowing developers to focus solely on building and deploying their applications. Azure Container Apps handles automatic scaling, and load balancing, and natively integrates with other Azure services, such as [Azure Monitor](#) and [Azure Container Registry \(ACR\)](#), to provide a comprehensive and secure application deployment experience.

Azure Container Apps offers benefits such as rapid deployment, easy scalability, cost-efficiency, and seamless integration with other Azure services, making it an attractive choice for modern application development and deployment scenarios.

Azure OpenAI Service

The [Azure OpenAI Service](#) is a platform offered by Microsoft Azure that provides cognitive services powered by [OpenAI](#) models. One of the models available through this service is the [ChatGPT](#) model, which is designed for interactive conversational tasks. It allows developers to integrate natural language understanding and generation capabilities into their applications.

Azure OpenAI Service provides REST API access to OpenAI's powerful language models including the [GPT-3](#), [Codex](#) and [Embeddings](#) model series. In addition, the new [GPT-4](#) and [ChatGPT](#) model series have now reached general availability. These models can be easily adapted to your specific task, including but not limited to content generation, summarization, semantic search, and natural language-to-code translation. Users can access the service through REST APIs, Python SDK, or our web-based interface in the Azure OpenAI Studio. You can use [Embeddings model](#) to transform raw data or inputs into meaningful and compact numerical representations called embeddings. Embeddings capture the semantic or contextual information of the input data in a lower-dimensional space, making it easier for machine learning algorithms to process and analyze the data effectively. Embeddings can be stored in a vector database, such as [ChromaDB](#) or [Facebook AI Similarity Search \(FAISS\)](#), explicitly designed for efficient storage, indexing, and retrieval of vector embeddings.

The [Chat Completion API](#), which is part of the Azure OpenAI Service, provides a dedicated interface for interacting with the [ChatGPT](#) and [GPT-4 models](#). This API is currently in preview and is the preferred method for accessing these models. The GPT-4 models can only be accessed through this API.

[GPT-3](#), [GPT-3.5](#), and [GPT-4](#) models from OpenAI are prompt-based. With prompt-based models, the user interacts with the model by entering a text prompt, to which the model responds with a text completion. This completion is the model's continuation of the input text. While these models are compelling, their behavior is also very sensitive to the prompt. This makes prompt construction a critical skill to develop. For more information, see [Introduction to prompt engineering](#).

Prompt construction can be complex. In practice, the prompt acts to configure the model weights to complete the desired task, but it's more of an art than a science, often requiring experience and intuition to craft a successful prompt. The goal of this article is to help get you started with this learning process. It attempts to capture general concepts and patterns that apply to all GPT models. However, it's essential to understand that each model behaves differently, so the learnings may not apply equally to all models.

Prompt engineering refers to the process of creating instructions called prompts for Large Language Models (LLMs), such as OpenAI's ChatGPT. With the immense potential of LLMs to solve a wide range of tasks, leveraging prompt engineering can empower us to save significant time and facilitate the development of impressive applications. It holds the key to unleashing the full capabilities of these huge models, transforming how we interact and benefit from them. For more information, see [Prompt engineering techniques](#).

Vector Databases

A [vector database](#) is a specialized database that goes beyond traditional storage by organizing information to simplify the search for similar items. Instead of merely storing words or numbers, it leverages vector embeddings - unique numerical representations of data. These embeddings capture meaning, context, and relationships. For instance, words are represented as vectors, whereas similar words have similar vector values. The applications of vector databases are numerous and powerful. In language processing, they facilitate the discovery of related documents or sentences. By comparing the vector embeddings of different texts, finding similar or related information becomes faster and more efficient. This capability benefits search engines and recommendation systems, which can suggest relevant articles or products based on user interests.

In the realm of image analysis, vector databases excel in finding visually similar images. By representing images as vectors, a simple comparison of vector values can identify visually similar images. This capability is precious for tasks like reverse image search or content-based image retrieval.

Additionally, vector databases find applications in fraud detection, anomaly detection, and clustering. By comparing vector embeddings of data points, unusual patterns can be detected, and similar items can be grouped together, aiding in effective data analysis and decision-making.

Here is a list of the most popular vector databases:

- [ChromaDB](#) is a powerful database solution that stores and retrieves vector embeddings efficiently. It is commonly used in AI applications, including chatbots and document analysis systems. By storing embeddings in ChromaDB, users can easily search and retrieve similar vectors, enabling faster and more accurate matching or recommendation processes. ChromaDB offers excellent scalability high performance, and supports various indexing techniques to optimize search operations. It is a versatile tool that enhances the functionality and efficiency of AI applications that rely on vector embeddings.
- [Facebook AI Similarity Search \(FAISS\)](#) is another widely used vector database. Facebook AI Research develops it and offers highly optimized algorithms for similarity search and clustering of vector embeddings. FAISS is known for its speed and scalability, making it suitable for large-scale applications. It offers different indexing methods like flat, IVF (Inverted File System), and HNSW (Hierarchical Navigable Small World) to organize and search vector data efficiently.
- [SingleStore](#): SingleStore aims to deliver the world's fastest distributed SQL database for data-intensive applications: SingleStoreDB, which combines transactional + analytical workloads in a single platform.

- Astra DB: DataStax Astra DB is a cloud-native, multi-cloud, fully managed database-as-a-service based on Apache Cassandra, which aims to accelerate application development and reduce deployment time for applications from weeks to minutes.
- Milvus: Milvus is an open source vector database built to power embedding similarity search and AI applications. Milvus makes unstructured data search more accessible and provides a consistent user experience regardless of the deployment environment. Milvus 2.0 is a cloud-native vector database with storage and computation separated by design. All components in this refactored version of Milvus are stateless to enhance elasticity and flexibility.
- Qdrant: Qdrant is a vector similarity search engine and database for AI applications. Along with open-source, Qdrant is also available in the cloud. It provides a production-ready service with an API to store, search, and manage points—vectors with an additional payload. Qdrant is tailored to extended filtering support. It makes it useful for all sorts of neural network or semantic-based matching, faceted search, and other applications.
- Pinecone: Pinecone is a fully managed vector database that makes adding vector search to production applications accessible. It combines state-of-the-art vector search libraries, advanced features such as filtering, and distributed infrastructure to provide high performance and reliability at any scale.
- Vespa: Vespa is a platform for applications combining data and AI online. Building such applications on Vespa helps users avoid integration work to get features, and it can scale to support any amount of traffic and data. To deliver that, Vespa provides a broad range of query capabilities, a computation engine with support for modern machine-learned models, hands-off operability, data management, and application development support. It is free and open source to use under the Apache 2.0 license.
- Zilliz: Milvus is an open-source vector database, with over 18,409 stars on GitHub and 3.4 million+ downloads. Milvus supports billion-scale vector search and has over 1,000 enterprise users. Zilliz Cloud provides a fully-managed Milvus service made by the creators of Milvus. This helps to simplify the process of deploying and scaling vector search applications by eliminating the need to create and maintain complex data infrastructure. As a DBaaS, Zilliz simplifies the process of deploying and scaling vector search applications by eliminating the need to create and maintain complex data infrastructure.
- Weaviate: Weaviate is an open-source vector database used to store data objects and vector embeddings from ML-models, and scale into billions of data objects from the same name company in Amsterdam. Users can index billions of data objects to search through and combine multiple search techniques, such as keyword-based and vector search, to provide search experiences.

This sample makes use of ChromaDB vector database, but you can easily modify the code to use another vector database. You can even use Azure Cache for Redis Enterprise to store the vector embeddings and compute vector similarity with high performance and low latency. For more information, see Vector Similarity Search with Azure Cache for Redis Enterprise

LangChain

LangChain is a software framework designed to streamline the development of applications using large language models (LLMs). It serves as a language model integration framework, facilitating various applications like document analysis and summarization, chatbots, and code analysis.

LangChain's integrations cover an extensive range of systems, tools, and services, making it a comprehensive solution for language model-based applications. LangChain integrates with the major cloud platforms such as Microsoft Azure, Amazon AWS, and Google, and with API wrappers for various purposes like news, movie information, and weather, as well as support for Bash, web scraping, and more. It also supports multiple

language models, including those from OpenAI, Anthropic, and Hugging Face. Moreover, LangChain offers various functionalities for document handling, code generation, analysis, debugging, and interaction with databases and other data sources.

Chainlit

[Chainlit](#) is an open-source Python package that is specifically designed to create user interfaces (UIs) for AI applications. It simplifies the process of building interactive chats and interfaces, making developing AI-powered applications faster and more efficient. While Streamlit is a general-purpose UI library, Chainlit is purpose-built for AI applications and seamlessly integrates with other AI technologies such as [LangChain](#), [LlamaIndex](#), and [LangFlow](#).

With Chainlit, developers can easily create intuitive UIs for their AI models, including ChatGPT-like applications. It provides a user-friendly interface for users to interact with AI models, enabling conversational experiences and information retrieval. Chainlit also offers unique features, such as displaying the [Chain of Thought](#), which allows users to explore the reasoning process directly within the UI. This feature enhances transparency and enables users to understand how the AI arrives at its responses or recommendations.

For more information, see the following resources:

- [Documentation](#)
- [Examples](#)
- [API Reference](#)
- [Cookbook](#)

Deploy the Infrastructure

Before deploying the Terraform modules in the `terraform/infra` folder, specify a value for the following variables in the [terraform.tfvars](#) variable definitions file.

```
1 | name_prefix = "Blue"
2 | location    = "EastUS"
```

This is the definition of each variable:

- `prefix` : specifies a prefix for all the Azure resources.
- `location` : specifies the region (e.g., EastUS) where deploying the Azure resources.

NOTE: Make sure to select a region where [Azure OpenAI Service \(AOAI\)](#) supports both [GPT-3.5/GPT-4](#) models like `gpt-35-turbo-16k` and [Embeddings](#) models like `text-embedding-ada-002`.

OpenAI Module

The following table contains the code from the `terraform/infra/modules/openai/main.tf` Terraform module used to deploy the [Azure OpenAI Service](#).

```

1 resource "azurerm_cognitive_account" "openai" {
2   name                = var.name
3   location             = var.location
4   resource_group_name = var.resource_group_name
5   kind                = "OpenAI"
6   custom_subdomain_name = var.custom_subdomain_name
7   sku_name             = var.sku_name
8   public_network_access_enabled = var.public_network_access_enabled
9   tags                = var.tags
10
11   identity {
12     type = "SystemAssigned"
13   }
14
15   lifecycle {
16     ignore_changes = [
17       tags
18     ]
19   }
20 }
21
22 resource "azurerm_cognitive_deployment" "deployment" {
23   for_each      = {for deployment in var.deployments: deployment.name => deployment}
24
25   name          = each.key
26   cognitive_account_id = azurerm_cognitive_account.openai.id
27
28   model {
29     format = "OpenAI"
30     name   = each.value.model.name
31     version = each.value.model.version
32   }
33
34   scale {
35     type = "Standard"
36   }
37 }
38
39 resource "azurerm_monitor_diagnostic_setting" "settings" {
40   name                = "DiagnosticsSettings"
41   target_resource_id  = azurerm_cognitive_account.openai.id
42   log_analytics_workspace_id = var.log_analytics_workspace_id
43
44   enabled_log {
45     category = "Audit"
46
47     retention_policy {
48       enabled = true
49       days    = var.log_analytics_retention_days
50     }
51   }
52
53   enabled_log {
54     category = "RequestResponse"
55
56     retention_policy {
57       enabled = true
58       days    = var.log_analytics_retention_days
59     }
60   }

```



```

61 }
62
63 enabled_log {
64     category = "Trace"
65
66     retention_policy {
67         enabled = true
68         days    = var.log_analytics_retention_days
69     }
70 }
71
72 metric {
73     category = "AllMetrics"
74
75     retention_policy {
76         enabled = true
77         days    = var.log_analytics_retention_days
78     }
79 }

```

Azure Cognitive Services uses custom subdomain names for each resource created through the [Azure portal](#), [Azure Cloud Shell](#), [Azure CLI](#), [Bicep](#), [Azure Resource Manager \(ARM\)](#), or [Terraform](#). Unlike regional endpoints, which were common for all customers in a specific Azure region, custom subdomain names are unique to the resource. Custom subdomain names are required to enable authentication features like Azure Active Directory (Azure AD). We need to specify a custom subdomain for our [Azure OpenAI Service](#), as our chatbot applications will use an Azure AD security token to access it. By default, the `terraform/infra/modules/openai/main.tf` module sets the value of the `custom_subdomain_name` parameter to the lowercase name of the Azure OpenAI resource. For more information on custom subdomains, see [Custom subdomain names for Cognitive Services](#).

This Terraform module allows you to pass an array containing the definition of one or more model deployments in the `deployments` variable. For more information on model deployments, see [Create a resource and deploy a model using Azure OpenAI](#). The `openai_deployments` variable in the `terraform/infra/variables.tf` file defines the structure and the default models deployed by the sample:

```

1  variable "openai_deployments" {
2      description = "(Optional) Specifies the deployments of the Azure OpenAI Service"
3      type = list(object({
4          name = string
5          model = object({
6              name = string
7              version = string
8          })
9          rai_policy_name = string
10     }))
11     default = [
12         {
13             name = "gpt-35-turbo-16k"
14             model = {
15                 name = "gpt-35-turbo-16k"
16                 version = "0613"
17             }
18             rai_policy_name = ""
19         },
20         {
21             name = "text-embedding-ada-002"
22             model = {
23                 name = "text-embedding-ada-002"
24                 version = "2"
25             }
26             rai_policy_name = ""
27         }
28     ]
29 }

```

Alternatively, you can use the [Terraform module for deploying Azure OpenAI Service](#) to deploy [Azure OpenAI Service](#).

Private Endpoint Module

The `terraform/infra/main.tf` the module creates [Azure Private Endpoints](#) and [Azure Private DNS Zones](#) for each of the following resources:

- [Azure OpenAI Service \(AOAI\)](#)
- [Azure Container Registry \(ACR\)](#)

In particular, it creates an [Azure Private Endpoint](#) and [Azure Private DNS Zone](#) to the [Azure OpenAI Service](#) as shown in the following code snippet:

```

1 module "openai_private_dns_zone" {
2     source                = "./modules/private_dns_zone"
3     name                  = "privatelink.openai.azure.com"
4     resource_group_name   = azurerm_resource_group.rg.name
5     tags                  = var.tags
6     virtual_networks_to_link = {
7         (module.virtual_network.name) = {
8             subscription_id = data.azurerm_client_config.current.subscription_id
9             resource_group_name = azurerm_resource_group.rg.name
10        }
11    }
12 }
13
14 module "openai_private_endpoint" {
15     source                = "./modules/private_endpoint"
16     name                  = "${module.openai.name}PrivateEndpoint"
17     location              = var.location
18     resource_group_name   = azurerm_resource_group.rg.name
19     subnet_id             = module.virtual_network.subnet_ids[var.vm_subnet_name]
20     tags                  = var.tags
21     private_connection_resource_id = module.openai.id
22     is_manual_connection  = false
23     subresource_name       = "account"
24     private_dns_zone_group_name = "AcrPrivateDnsZoneGroup"
25     private_dns_zone_group_ids = [module.openai_private_dns_zone.id]
26 }

```

Below you can read the code of the `terraform/infra/modules/private_endpoint/main.tf` module, which is used to create [Azure Private Endpoints](#):

```

1 resource "azurerm_private_endpoint" "private_endpoint" {
2     name                = var.name
3     location            = var.location
4     resource_group_name = var.resource_group_name
5     subnet_id          = var.subnet_id
6     tags               = var.tags
7
8     private_service_connection {
9         name                = "${var.name}Connection"
10        private_connection_resource_id = var.private_connection_resource_id
11        is_manual_connection = var.is_manual_connection
12        subresource_names    = try([var.subresource_name], null)
13        request_message      = try(var.request_message, null)
14    }
15
16    private_dns_zone_group {
17        name                = var.private_dns_zone_group_name
18        private_dns_zone_ids = var.private_dns_zone_group_ids
19    }
20
21    lifecycle {
22        ignore_changes = [
23            tags
24        ]
25    }
26 }

```

Private DNS Zone Module

In the following box, you can read the code of

the `terraform/infra/modules/private_dns_zone/main.tf` module, which is utilized to create the [Azure Private DNS Zones](#).

```

1 resource "azurerm_private_dns_zone" "private_dns_zone" {
2     name                = var.name
3     resource_group_name = var.resource_group_name
4     tags                = var.tags
5
6     lifecycle {
7         ignore_changes = [
8             tags
9         ]
10    }
11 }
12
13 resource "azurerm_private_dns_zone_virtual_network_link" "link" {
14     for_each = var.virtual_networks_to_link
15
16     name                = "link_to_${lower(basename(each.key))}"
17     resource_group_name = var.resource_group_name
18     private_dns_zone_name = azurerm_private_dns_zone.private_dns_zone.name
19     virtual_network_id   = "/subscriptions/${each.value.subscription_id}/resourceGroups/${each.value.resource_group_name}/virtualNetworks/${each.value.virtual_network_name}"
20
21     lifecycle {
22         ignore_changes = [
23             tags
24         ]
25    }
26 }

```

Workload Managed Identity Module

Below you can read the code of the `terraform/infra/modules/managed_identity/main.tf` module, which is used to create the [Azure Managed Identity](#) used by the [Azure Container Apps](#) to pull container images from the [Azure Container Registry](#), and by the chat applications to connect to the [Azure OpenAI Service](#). You can use a system-assigned or user-assigned managed identity from [Azure Active Directory \(Azure AD\)](#) to let [Azure Container Apps](#) access any Azure AD-protected resource. For more information, see [Managed identities in Azure Container Apps](#). You can pull container images from private repositories in an [Azure Container Registry](#) using user-assigned or user-assigned managed identities for authentication to avoid using administrative credentials. For more information, see [Azure Container Apps image pull with managed identity](#). This user-defined managed identity is assigned the [Cognitive Services User](#) role on the [Azure OpenAI Service](#) namespace and [ACRPull](#) role on the [Azure Container Registry \(ACR\)](#). By assigning the above roles, you grant the user-defined managed identity access to these resources.

```

1 resource "azurerm_user_assigned_identity" "workload_user_assigned_identity" {
2     name                        = var.name
3     resource_group_name = var.resource_group_name
4     location              = var.location
5     tags                   = var.tags
6
7     lifecycle {
8         ignore_changes = [
9             tags
10        ]
11    }
12 }
13
14 resource "azurerm_role_assignment" "cognitive_services_user_assignment" {
15     scope                        = var.openai_id
16     role_definition_name = "Cognitive Services User"
17     principal_id              = azurerm_user_assigned_identity.workload_user_assigned_identity.principal_id
18     skip_service_principal_aad_check = true
19 }
20
21 resource "azurerm_role_assignment" "acr_pull_assignment" {
22     scope                        = var.acr_id
23     role_definition_name = "AcrPull"
24     principal_id              = azurerm_user_assigned_identity.workload_user_assigned_identity.principal_id
25     skip_service_principal_aad_check = true
26 }

```

Deploy the Applications

Before deploying the Terraform modules in the `terraform/apps` folder, specify a value for the following variables in the `Terraform.tfvars` variable definitions file.

```

1 resource_group_name      = "BlueRG"
2 container_app_environment_name = "BlueEnvironment"
3 container_registry_name    = "BlueRegistry"
4 workload_managed_identity_name = "BlueWorkloadIdentity"
5 container_apps            = [
6     {
7         name                = "chatapp"
8         revision_mode        = "Single"
9         ingress              = {
10             allow_insecure_connections = true
11             external_enabled           = true
12             target_port                = 8000
13             transport                 = "http"
14             traffic_weight             = {
15                 label                = "default"
16                 latest_revision       = true
17                 revision_suffix       = "default"
18                 percentage            = 100
19             }
20         }
21         template             = {
22             containers        = [
23                 {
24                     name        = "chat"
25                     image       = "chat:v1"
26                     cpu          = 0.5
27                     memory       = "1Gi"
28                     env          = [
29                         {
30                             name    = "TEMPERATURE"
31                             value   = 0.9
32                         },
33                         {
34                             name    = "AZURE_OPENAI_BASE"
35                             value   = "https://blueopenai.openai.azure.com/"
36                         },
37                         {
38                             name    = "AZURE_OPENAI_KEY"
39                             value   = ""
40                         },
41                         {
42                             name    = "AZURE_OPENAI_TYPE"
43                             value   = "azure_ad"
44                         },
45                         {
46                             name    = "AZURE_OPENAI_VERSION"
47                             value   = "2023-06-01-preview"
48                         },
49                         {
50                             name    = "AZURE_OPENAI_DEPLOYMENT"
51                             value   = "gpt-35-turbo-16k"
52                         },
53                         {
54                             name    = "AZURE_OPENAI_MODEL"
55                             value   = "gpt-35-turbo-16k"
56                         },
57                         {
58                             name    = "AZURE_OPENAI_SYSTEM_MESSAGE"
59                             value   = "You are a helpful assistant."
60

```

```

61         },
62         {
63             name          = "MAX_RETRIES"
64             value          = 5
65         },
66         {
67             name          = "BACKOFF_IN_SECONDS"
68             value          = "1"
69         },
70         {
71             name          = "TOKEN_REFRESH_INTERVAL"
72             value          = 2700
73         }
74     ]
75     liveness_probe        = {
76         failure_count_threshold = 3
77         initial_delay          = 30
78         interval_seconds       = 60
79         path                   = "/"
80         port                   = 8000
81         timeout                = 30
82         transport              = "HTTP"
83     }
84     readiness_probe = {
85         failure_count_threshold = 3
86         interval_seconds       = 60
87         path                   = "/"
88         port                   = 8000
89         success_count_threshold = 3
90         timeout                = 30
91         transport              = "HTTP"
92     }
93     startup_probe = {
94         failure_count_threshold = 3
95         interval_seconds       = 60
96         path                   = "/"
97         port                   = 8000
98         timeout                = 30
99         transport              = "HTTP"
100    }
101    }
102    ]
103    min_replicas          = 1
104    max_replicas          = 3
105    }
106    },
107    {
108        name              = "docapp"
109        revision_mode      = "Single"
110        ingress            = {
111            allow_insecure_connections = true
112            external_enabled           = true
113            target_port                = 8000
114            transport                  = "http"
115            traffic_weight              = {
116                label                  = "default"
117                latest_revision        = true
118                revision_suffix        = "default"
119                percentage              = 100
120            }

```



```

121     }
122     template                                     = {
123     containers                                 = [
124         {
125             name                               = "doc"
126             image                             = "doc:v1"
127             cpu                               = 0.5
128             memory                           = "1Gi"
129             env                               = [
130                 {
131                     name                       = "TEMPERATURE"
132                     value                     = 0.9
133                 },
134                 {
135                     name                       = "AZURE_OPENAI_BASE"
136                     value                     = "https://blueopenai.openai.azure.com/"
137                 },
138                 {
139                     name                       = "AZURE_OPENAI_KEY"
140                     value                     = ""
141                 },
142                 {
143                     name                       = "AZURE_OPENAI_TYPE"
144                     value                     = "azure_ad"
145                 },
146                 {
147                     name                       = "AZURE_OPENAI_VERSION"
148                     value                     = "2023-06-01-preview"
149                 },
150                 {
151                     name                       = "AZURE_OPENAI_DEPLOYMENT"
152                     value                     = "gpt-35-turbo-16k"
153                 },
154                 {
155                     name                       = "AZURE_OPENAI_MODEL"
156                     value                     = "gpt-35-turbo-16k"
157                 },
158                 {
159                     name                       = "AZURE_OPENAI_ADA_DEPLOYMENT"
160                     value                     = "text-embedding-ada-002"
161                 },
162                 {
163                     name                       = "AZURE_OPENAI_SYSTEM_MESSAGE"
164                     value                     = "You are a helpful assistant."
165                 },
166                 {
167                     name                       = "MAX_RETRIES"
168                     value                     = 5
169                 },
170                 {
171                     name                       = "CHAINLIT_MAX_FILES"
172                     value                     = 10
173                 },
174                 {
175                     name                       = "TEXT_SPLITTER_CHUNK_SIZE"
176                     value                     = 1000
177                 },
178                 {
179                     name                       = "TEXT_SPLITTER_CHUNK_OVERLAP"
180                     value                     = 100
181

```

```

182     },
183     {
184         name          = "EMBEDDINGS_CHUNK_SIZE"
185         value          = 16
186     },
187     {
188         name          = "BACKOFF_IN_SECONDS"
189         value          = "1"
190     },
191     {
192         name          = "CHAINLIT_MAX_SIZE_MB"
193         value          = 100
194     },
195     {
196         name          = "TOKEN_REFRESH_INTERVAL"
197         value          = 2700
198     }
199 ]
200 liveness_probe = {
201     failure_count_threshold = 3
202     initial_delay           = 30
203     interval_seconds        = 60
204     path                    = "/"
205     port                    = 8000
206     timeout                 = 30
207     transport               = "HTTP"
208 }
209 readiness_probe = {
210     failure_count_threshold = 3
211     interval_seconds        = 60
212     path                    = "/"
213     port                    = 8000
214     success_count_threshold = 3
215     timeout                 = 30
216     transport               = "HTTP"
217 }
218 startup_probe = {
219     failure_count_threshold = 3
220     interval_seconds        = 60
221     path                    = "/"
222     port                    = 8000
223     timeout                 = 30
224     transport               = "HTTP"
225 }
226 }
227 ]
228 min_replicas      = 1
229 max_replicas      = 3
    }
}]

```

This is the definition of each variable:

- `resource_group_name` : specifies the name of the resource group that contains the infrastructure resources: [Azure OpenAI Service](#), [Azure Container Registry](#), [Azure Container Apps Environment](#), [Azure Log Analytics](#), and [user-defined managed identity](#).
- `container_app_environment_name` : the name of the [Azure Container Apps Environment](#) in which to deploy the chat applications. [Skip to final configuration](#)

- `container_registry_name` : the name of [Azure Container Registry](#) used to hold the container images of the chat applications.
- `workload_managed_identity_name` : the name of the [user-defined managed identity](#) used by the chat applications to authenticate with [Azure OpenAI Service](#) and [Azure Container Registry](#).
- `container_apps` : the definition of the two chat applications. The application configuration does not specify the following data because the `container_app` module later defines this information:
 - `image` : This field contains the name and tag of the container image but not the login server of the [Azure Container Registry](#).
 - `identity` : The identity of the container app.
 - `registry` : The registry hosting the container image for the application.
 - `AZURE_CLIENT_ID` : The client ID of the user-defined managed identity used by the application to authenticate with [Azure OpenAI Service](#) and [Azure Container Registry](#).
 - `AZURE_OPENAI_TYPE` : This environment variable specifies the authentication type with [Azure OpenAI Service](#): if you set the value of the `AZURE_OPENAI_TYPE` environment variable to `azure`, you need to specify the OpenAI key as a value for the `AZURE_OPENAI_KEY` environment variable. Instead, if you set the value to `azure_ad` in the application code, assign an Azure AD security token to the `openai_api_key` property. For more information, see [How to switch between OpenAI and Azure OpenAI endpoints with Python](#).

Container App Module

The `terraform/apps/modules/container_app/main.tf` module is utilized to create the [Azure Container Apps](#). The module defines and uses the following [data source](#) for the [Azure Container Registry](#), [Azure Container Apps Environment](#), and [user-defined managed identity](#) created when deploying the infrastructure. These data sources are used to access the properties of these Azure resources.

```

1 data "azurerm_container_app_environment" "container_app_environment" {
2   name                = var.container_app_environment_name
3   resource_group_name = var.resource_group_name
4 }
5
6 data "azurerm_container_registry" "container_registry" {
7   name                = var.container_registry_name
8   resource_group_name = var.resource_group_name
9 }
10
11 data "azurerm_user_assigned_identity" "workload_user_assigned_identity" {
12   name                = var.workload_managed_identity_name
13   resource_group_name = var.resource_group_name
14 }

```

The module creates and utilizes the following local variables:

```

1 | locals {
2 |     identity = {
3 |         type      = "UserAssigned"
4 |         identity_ids = [data.azurerm_user_assigned_identity.workload_user_assigned_identity.id]
5 |     }
6 |     identity_env = {
7 |         name      = "AZURE_CLIENT_ID"
8 |         secret_name = null
9 |         value      = data.azurerm_user_assigned_identity.workload_user_assigned_identity.client_id
10 |    }
11 |    registry = {
12 |        server      = data.azurerm_container_registry.container_registry.login_server
13 |        identity     = data.azurerm_user_assigned_identity.workload_user_assigned_identity.id
14 |    }
15 | }

```

This is the explanation of each local variable:

- `identity` : uses the resource ID of the user-defined managed identity to define the `identity` block for each container app deployed by the module.
- `identity_env` : uses the client ID of the user-defined managed identity to define the value of the `AZURE_CLIENT_ID` environment variable that is appended to the list of environment variables of each container app deployed by the module.
- `registry` : uses the login server of the Azure Container Registry to define the `registry` block for each container app deployed by the module.

Here is the complete Terraform code of the module:

```

data "azurerm_container_app_environment" "container_app_environment" {
  name                = var.container_app_environment_name
  resource_group_name = var.resource_group_name
}

data "azurerm_container_registry" "container_registry" {
  name                = var.container_registry_name
  resource_group_name = var.resource_group_name
}

data "azurerm_user_assigned_identity" "workload_user_assigned_identity" {
  name                = var.workload_managed_identity_name
  resource_group_name = var.resource_group_name
}

locals {
  identity = {
    type           = "UserAssigned"
    identity_ids   = [data.azurerm_user_assigned_identity.workload_user_assigned_identity.id]
  }
  identity_env = {
    name          = "AZURE_CLIENT_ID"
    secret_name    = null
    value         = data.azurerm_user_assigned_identity.workload_user_assigned_identity.client_id
  }
  registry = {
    server        = data.azurerm_container_registry.container_registry.login_server
    identity       = data.azurerm_user_assigned_identity.workload_user_assigned_identity.id
  }
}

resource "azurerm_container_app" "container_app" {
  for_each          = {for app in var.container_apps: app.name => app}

  container_app_environment_id = data.azurerm_container_app_environment.container_app_environment
  name                        = each.key
  resource_group_name        = var.resource_group_name
  revision_mode              = each.value.revision_mode
  tags                       = each.value.tags

  template {
    max_replicas    = each.value.template.max_replicas
    min_replicas    = each.value.template.min_replicas
    revision_suffix = each.value.template.revision_suffix

    dynamic "container" {
      for_each = each.value.template.containers

      content {
        cpu       = container.value.cpu
        image     = "${data.azurerm_container_registry.container_registry.login_server}/${container
        memory    = container.value.memory
        name      = container.value.name
        args      = container.value.args
        command   = container.value.command

        dynamic "env" {
          for_each = container.value.env == null ? [local.identity_env] : concat(container.value.

```

```

    content {
      name      = env.value.name
      secret_name = env.value.secret_name
      value      = env.value.value
    }
  }

dynamic "liveness_probe" {
  for_each = container.value.liveness_probe == null ? [] : [container.value.liveness_probe]

  content {
    port                = liveness_probe.value.port
    transport            = liveness_probe.value.transport
    failure_count_threshold = liveness_probe.value.failure_count_threshold
    host                = liveness_probe.value.host
    initial_delay        = liveness_probe.value.initial_delay
    interval_seconds      = liveness_probe.value.interval_seconds
    path                 = liveness_probe.value.path
    timeout              = liveness_probe.value.timeout

    dynamic "header" {
      for_each = liveness_probe.value.header == null ? [] : [liveness_probe.value.header]

      content {
        name = header.value.name
        value = header.value.value
      }
    }
  }
}

dynamic "readiness_probe" {
  for_each = container.value.readiness_probe == null ? [] : [container.value.readiness_probe]

  content {
    port                = readiness_probe.value.port
    transport            = readiness_probe.value.transport
    failure_count_threshold = readiness_probe.value.failure_count_threshold
    host                = readiness_probe.value.host
    interval_seconds      = readiness_probe.value.interval_seconds
    path                 = readiness_probe.value.path
    success_count_threshold = readiness_probe.value.success_count_threshold
    timeout              = readiness_probe.value.timeout

    dynamic "header" {
      for_each = readiness_probe.value.header == null ? [] : [readiness_probe.value.header]

      content {
        name = header.value.name
        value = header.value.value
      }
    }
  }
}

dynamic "startup_probe" {
  for_each = container.value.startup_probe == null ? [] : [container.value.startup_probe]

  content {
    port                = startup_probe.value.port
    transport            = startup_probe.value.transport
    failure_count_threshold = startup_probe.value.failure_count_threshold
    host                = startup_probe.value.host
    interval_seconds      = startup_probe.value.interval_seconds
    path                 = startup_probe.value.path
    success_count_threshold = startup_probe.value.success_count_threshold
    timeout              = startup_probe.value.timeout

    dynamic "header" {
      for_each = startup_probe.value.header == null ? [] : [startup_probe.value.header]

      content {
        name = header.value.name
        value = header.value.value
      }
    }
  }
}

```

```

transport          = startup_probe.value.transport
failure_count_threshold = startup_probe.value.failure_count_threshold
host               = startup_probe.value.host
interval_seconds   = startup_probe.value.interval_seconds
path               = startup_probe.value.path
timeout            = startup_probe.value.timeout

dynamic "header" {
  for_each = startup_probe.value.header == null ? [] : [startup_probe.value.header]

  content {
    name = header.value.name
    value = header.value.name
  }
}

dynamic "volume_mounts" {
  for_each = container.value.volume_mounts == null ? [] : [container.value.volume_mounts]

  content {
    name = volume_mounts.value.name
    path = volume_mounts.value.path
  }
}

dynamic "volume" {
  for_each = each.value.template.volume == null ? [] : each.value.template.volume

  content {
    name          = volume.value.name
    storage_name = volume.value.storage_name
    storage_type  = volume.value.storage_type
  }
}

dynamic "dapr" {
  for_each = each.value.dapr == null ? [] : [each.value.dapr]

  content {
    app_id      = dapr.value.app_id
    app_port    = dapr.value.app_port
    app_protocol = dapr.value.app_protocol
  }
}

dynamic "identity" {
  for_each = each.value.identity == null ? [local.identity] : [each.value.identity]

  content {
    type          = identity.value.type
    identity_ids = identity.value.identity_ids
  }
}

dynamic "ingress" {
  for_each = each.value.ingress == null ? [] : [each.value.ingress]

  content {
    name = ingress.value.name
    path = ingress.value.path
    port = ingress.value.port
  }
}

```

```

182     for_each = each.value.ingress == null ? [] : [each.value.ingress]
183
184     content {
185         target_port          = ingress.value.target_port
186         allow_insecure_connections = ingress.value.allow_insecure_connections
187         external_enabled      = ingress.value.external_enabled
188         transport             = ingress.value.transport
189
190         dynamic "traffic_weight" {
191             for_each = ingress.value.traffic_weight == null ? [] : [ingress.value.traffic_weight]
192
193             content {
194                 percentage      = traffic_weight.value.percentage
195                 label           = traffic_weight.value.label
196                 latest_revision = traffic_weight.value.latest_revision
197                 revision_suffix = traffic_weight.value.revision_suffix
198             }
199         }
200     }
201 }
202
203 dynamic "registry" {
204     for_each = each.value.registry == null ? [local.registry] : concat(each.value.registry, [local.registry])
205
206     content {
207         server = registry.value.server
208         identity = registry.value.identity
209     }
210 }
211
212 dynamic "secret" {
213     for_each = nonsensitive(toiset([for pair in lookup(var.container_app_secrets, each.key, []) :
214
215         content {
216             name = secret.key
217             value = local.container_app_secrets[each.key][secret.key]
218         }
219     ]))
220 }

```

As you can notice, the module uses the login server of the [Azure Container Registry](#) to create the fully qualified name of the container image of the current container app.

Managed identities in Azure Container Apps

Each chat application makes use of a [DefaultAzureCredential](#) object to acquire a security token from Azure Active Directory and authenticate and authorize with [Azure OpenAI Service \(AOAI\)](#) and [Azure Container Registry \(ACR\)](#) using the credentials of the user-defined managed identity associated with the container app. You can use a managed identity in a running container app to authenticate and authorize with any [service that supports Azure AD authentication](#). With managed identities:

- Container apps and applications connect to resources with the managed identity. You don't need to manage credentials in your container apps.
- You can use role-based access control to grant specific permissions to a managed identity.

- System-assigned identities are automatically created and managed. They are deleted when your container app or container app is deleted.
- You can add and delete user-assigned identities and assign them to multiple resources. They are independent of your container app or the container app's lifecycle.
- You can use managed identity to [authenticate with a private Azure Container Registry](#) without a username and password to pull containers for your Container App.
- You can use [managed identity to create connections for Dapr-enabled applications via Dapr components](#)

For more information, see [Managed identities in Azure Container Apps](#). The workloads running in a container app can use the Azure Identity client libraries to acquire a security token from the Azure Active Directory. You can choose one of the following approaches inside your code:

- Use `DefaultAzureCredential`, which will attempt to use the `WorkloadIdentityCredential`.
- Create a `ChainedTokenCredential` instance that includes `WorkloadIdentityCredential`.
- Use `WorkloadIdentityCredential` directly.

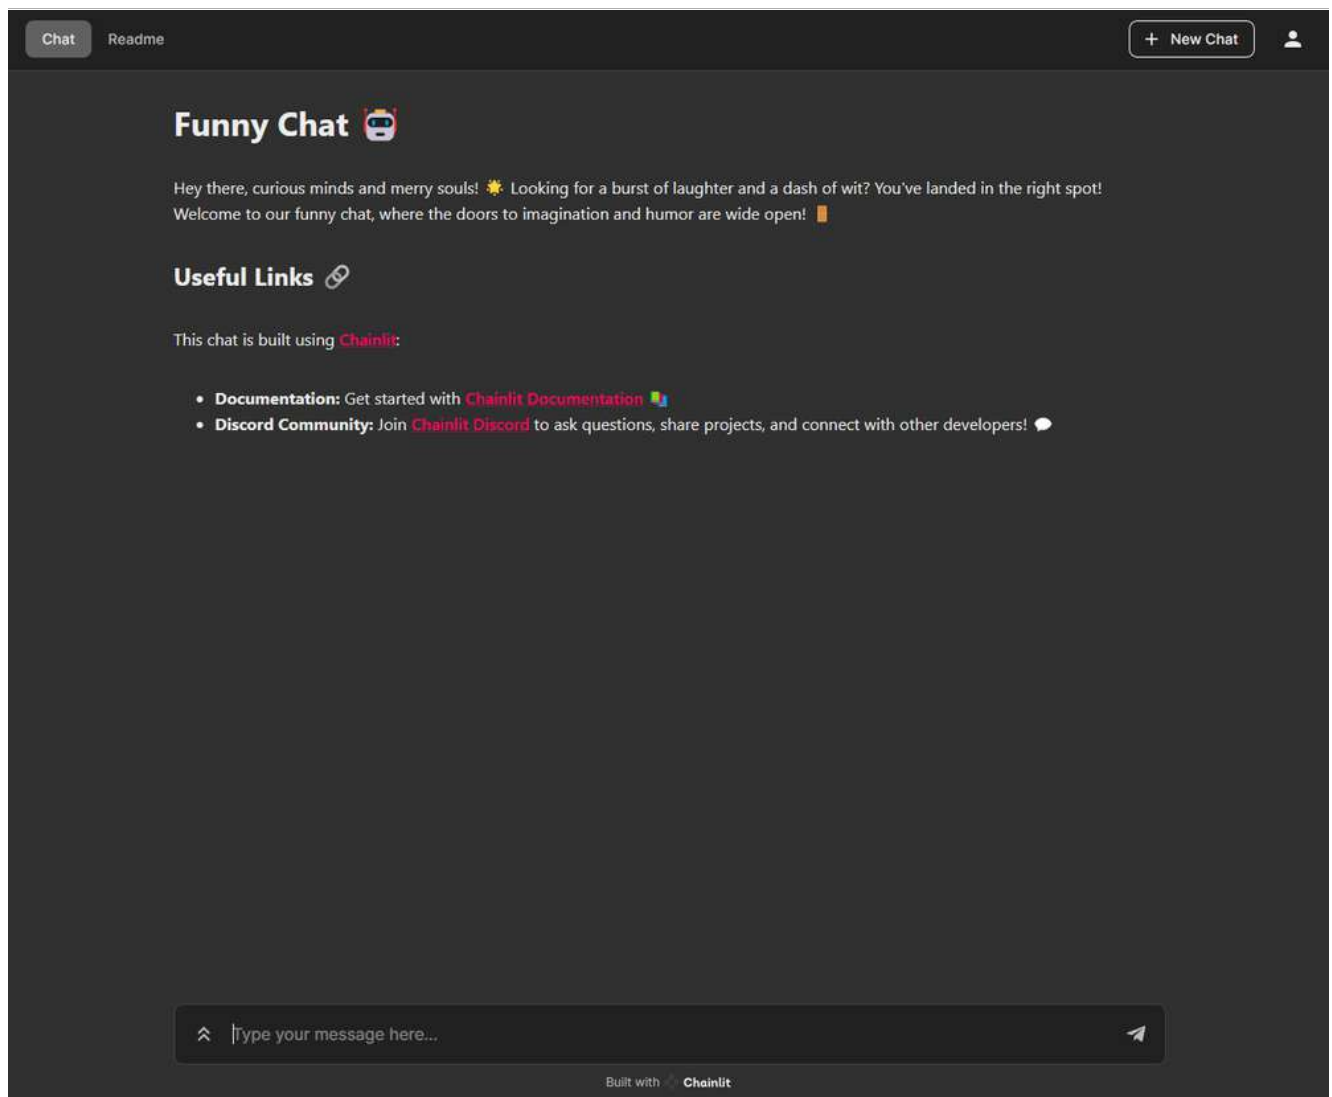
The following table provides the minimum package version required for each language's client library.

Language	Library	Minimum Version	Example
.NET	Azure.Identity	1.9.0	Link
Go	azidentity	1.3.0	Link
Java	azure-identity	1.9.0	Link
JavaScript	@azure/identity	3.2.0	Link
Python	azure-identity	1.13.0	Link

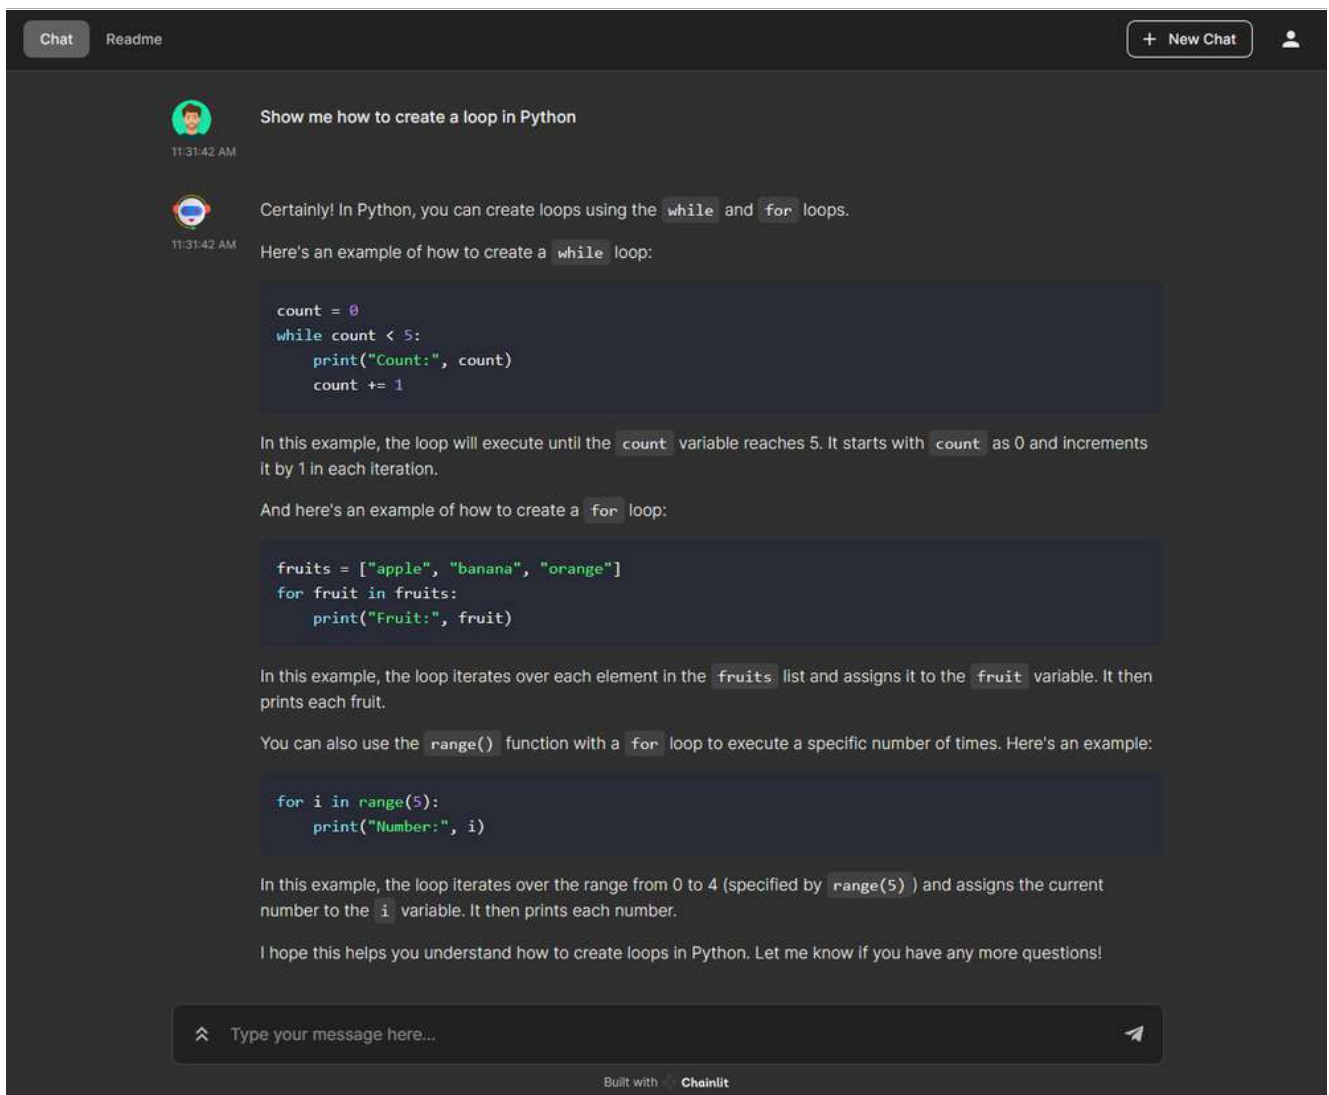
NOTE: When using Azure Identity client library with Azure Container Apps, the client ID of the managed identity must be specified. When using the `DefaultAzureCredential`, you can explicitly specify the client ID of the container app managed identity in the `AZURE_CLIENT_ID` environment variable.

Simple Chat Application

The Simple Chat Application is a large language model-based chatbot that allows users to submit general-purpose questions to a [GPT](#) model, generating and streaming back human-like and engaging conversational responses. The following picture shows the welcome screen of the chat application.



You can modify the welcome screen in markdown by editing the `chainlit.md` file at the project's root. If you do not want a welcome screen, leave the file empty. The following picture shows what happens when a user submits a new message in the chat.



Chainlit can render messages in markdown format and provides classes to support the following elements:

- **Audio:** The `Audio` class allows you to display an audio player for a specific audio file in the chatbot user interface. You must provide either a URL or a path or content bytes.
- **Avatar:** The `Avatar` class allows you to display an avatar image next to a message instead of the author's name. You need to send the element once. Next,, if an avatar's name matches an author's name, the avatar will be automatically displayed. You must provide either a URL or a path or content bytes.
- **File:** The `File` class allows you to display a button that lets users download the file's content. You must provide either a URL or a path or content bytes.
- **Image:** The `Image` class is designed to create and handle image elements to be sent and displayed in the chatbot user interface. You must provide either a URL or a path or content bytes.
- **Pdf:** The `pdf` class allows you to display a PDF hosted remotely or locally in the chatbot UI. This class either takes a URL of a PDF hosted online or the path of a local PDF.
- **Pyplot:** The `Pyplot` class allows you to display a `Matplotlib` pyplot chart in the chatbot UI. This class takes a pyplot figure.
- **TaskList:** The `TaskList` class allows you to display a task list next to the chatbot UI.
- **Text:** The `Text` class allows you to display a text element in the chatbot UI. This class takes a string and creates a text element that can be sent to the UI. It supports the markdown syntax for formatting text. You must provide either a URL or a path or content bytes.

Chainlit provides three [display options](#) that determine how an element is rendered in the context of its use. The `ElementDisplay` type represents these options. The following display options are available:

- `Side` : this option displays the element on a sidebar. The sidebar is hidden by default and opened upon element reference click.
- `Page` : this option displays the element on a separate page. The user is redirected to the page upon an element reference click.
- `Inline` : this option displays the element below the message. If the element is [global](#), it is displayed if it is explicitly mentioned in the message. If the element is [scoped](#), it is displayed regardless of whether it is expressly mentioned in the message.

You can click the user icon on the UI to access the chat settings and choose, for example, between the light and dark themes.

The application is built in Python. Let's take a look at the individual parts of the application code. The Python code starts by importing the necessary packages/modules in the following section.

```
1  # Import packages
2  import os
3  import sys
4  import time
5  import openai
6  import random
7  import logging
8  import chainlit as cl
9  from azure.identity import DefaultAzureCredential
10 from dotenv import load_dotenv
11 from dotenv import dotenv_values
12
13 # Load environment variables from .env file
14 if os.path.exists(".env"):
15     load_dotenv(override=True)
16     config = dotenv_values(".env")
```

These are the libraries used by the chat application:

1. `os` : This module provides a way of interacting with the operating system, enabling the code to access environment variables, file paths, etc.
2. `sys` : This module provides access to some variables used or maintained by the interpreter and functions that interact with the interpreter.
3. `time` : This module provides various time-related time manipulation and measurement functions.
4. `openai` : The OpenAI Python library provides convenient access to the OpenAI API from applications written in Python. It includes a pre-defined set of classes for API resources that initialize themselves dynamically from API responses, making it compatible with a wide range of versions of the OpenAI API. You can find usage examples for the OpenAI Python library in our [API reference](#) and the [OpenAI Cookbook](#).
5. `random` : This module provides functions to generate random numbers.
6. `logging` : This module provides flexible logging of messages.
7. `chainlit as cl` : This imports the [Chainlit](#) library and aliases it as `cl` . Chainlit is used to create the UI of the application.

[Skip to later navigation](#)

8. `DefaultAzureCredential` from `azure.identity`: when the `openai_type` property value is `azure_ad`, a `DefaultAzureCredential` object from the [Azure Identity client library for Python - version 1.13.0(<https://learn.microsoft.com/en-us/python/api/overview/azure/identity-readme?view=azure-python>)] is used to acquire security token from the Azure Active Directory using the credentials of the user-defined managed identity, whose client ID is defined in the `AZURE_CLIENT_ID` environment variable.
9. `load_dotenv` and `dotenv_values` from `dotenv`: [Python-dotenv](#) reads key-value pairs from a `.env` file and can set them as environment variables. It helps in the development of applications following the [12-factor](#) principles.

The `requirements.txt` file under the `src` folder contains the list of packages used by the chat applications. You can restore these packages in your environment using the following command:

```
1 | pip install -r requirements.txt --upgrade
```

Next, the code reads environment variables and configures the OpenAI settings.

```
1 | # Read environment variables
2 | temperature = float(os.environ.get("TEMPERATURE", 0.9))
3 | api_base = os.getenv("AZURE_OPENAI_BASE")
4 | api_key = os.getenv("AZURE_OPENAI_KEY")
5 | api_type = os.environ.get("AZURE_OPENAI_TYPE", "azure")
6 | api_version = os.environ.get("AZURE_OPENAI_VERSION", "2023-06-01-preview")
7 | engine = os.getenv("AZURE_OPENAI_DEPLOYMENT")
8 | model = os.getenv("AZURE_OPENAI_MODEL")
9 | system_content = os.getenv("AZURE_OPENAI_SYSTEM_MESSAGE", "You are a helpful assistant.")
10 | max_retries = int(os.getenv("MAX_RETRIES", 5))
11 | backoff_in_seconds = float(os.getenv("BACKOFF_IN_SECONDS", 1))
12 |
13 | # Configure OpenAI
14 | openai.api_type = api_type
15 | openai.api_version = api_version
16 | openai.api_base = api_base
17 | openai.api_key = api_key
```

Here's a brief explanation of each variable and related environment variable:

1. `temperature`: A float value representing the temperature for [Create chat completion](#) method of the OpenAI API. It is fetched from the environment variables with a default value of 0.9.
2. `api_base`: The base URL for the OpenAI API.
3. `api_key`: The API key for the OpenAI API.
4. `api_type`: A string representing the type of the OpenAI API.
5. `api_version`: A string representing the version of the OpenAI API.
6. `engine`: The engine used for OpenAI API calls.
7. `model`: The model used for OpenAI API calls.
8. `system_content`: The content of the system message used for OpenAI API calls.
9. `max_retries`: The maximum number of retries for OpenAI API calls.
10. `backoff_in_seconds`: The backoff time in seconds for retries in case of failures.

In the next section, the code sets the default Azure credential based on the `api_type` and configures a logger for logging purposes.

```
1 # Set default Azure credential
2 default_credential = DefaultAzureCredential() if openai.api_type == "azure_ad" else None
3
4 # Configure a logger
5 logging.basicConfig(
6     stream=sys.stdout,
7     format='[%(asctime)s] %(filename)s:%(lineno)d} %(levelname)s - %(message)s',
8     level=logging.INFO
9 )
10 logger = logging.getLogger(__name__)
```

Here's a brief explanation:

1. `default_credential` : It sets the default Azure credential to `DefaultAzureCredential()` if the `api_type` is "azure_ad"; otherwise, it is set to `None`.
2. `logging.basicConfig()` : This function configures the logging system with specific settings.
 - `stream` : The output stream where log messages will be written. Here, it is set to `sys.stdout` for writing log messages to the standard output.
 - `format` : The format string for log messages. It includes the timestamp, filename, line number, log level, and the actual log message.
 - `level` : The logging level. It is set to `logging.INFO`, meaning only messages with the level `INFO` and above will be logged.
3. `logger` : This creates a logger instance named after the current module (`__name__`). The logger will be used to log messages throughout the code.

Next, the code defines a helper function `backoff` that takes an integer `attempt` and returns a float value representing the backoff time for exponential retries in case of API call failures.

```
1 def backoff(attempt: int) -> float:
2     return backoff_in_seconds * 2 ** attempt + random.uniform(0, 1)
```

The backoff time is calculated using the `backoff_in_seconds` and `attempt` variables. It follows the formula `backoff_in_seconds * 2 ** attempt + random.uniform(0, 1)`. This formula increases the backoff time exponentially with each attempt and adds a random value between 0 and 1 to avoid synchronized retries. Then, the application defines a function called `refresh_openai_token()` to refresh the OpenAI security token if needed.

```
1 def refresh_openai_token():
2     token = cl.user_session.get('openai_token')
3     if token is None or token.expires_on < int(time.time()) - 1800:
4         cl.user_session.set('openai_token', default_credential.get_token("https://cognitiveservice
5         openai.api_key = cl.user_session.get('openai_token').token
```

The function follows these steps:

1. It fetches the current token from `cl.user_session` (which seems to be a part of the `chainlit` library) using the key `'openai_token'`. The `user_session` is a dictionary that stores the user's session data. The `id` and `env` keys are reserved for the session ID and environment variables. Other keys can be used to store arbitrary data in the user's session.

[Skip to final configuration](#)

2. It checks if the token is `None` or if its expiration time (`expires_on`) is less than the current time minus 1800 seconds (30 minutes).

Next, the code defines a function called `start_chat` that is used to initialize the when the user connects to the application or clicks the `New Chat` button.

```
1  .on_chat_start
2  async def start_chat():
3      # Sending Avatars for Chat Participants
4      await cl.Avatar(
5          name="Chatbot",
6          url="https://cdn-icons-png.flaticon.com/512/8649/8649595.png"
7      ).send()
8      await cl.Avatar(
9          name="Error",
10         url="https://cdn-icons-png.flaticon.com/512/8649/8649595.png"
11     ).send()
12     await cl.Avatar(
13         name="User",
14         url="https://media.architecturaldigest.com/photos/5f241de2c850b2a36b415024/master/w_1600%
15     ).send()
16
17     # Initializing message_history in user session
18     system_content = "Welcome to the chat!"
19     cl.user_session.set("message_history", [{"role": "system", "content": system_content}])
```

Here is a brief explanation of the function steps:

- `@cl.on_chat_start` : The `on_chat_start` decorator registers a callback function `start_chat()` to be called when the Chainlit chat starts. It is used to set up the chat and send avatars for the Chatbot, Error, and User participants in the chat.
- `cl.Avatar()` : the Avatar class allows you to display an avatar image next to a message instead of the author name. You need to send the element once. Next, if an avatar's name matches an author's name, the avatar will be automatically displayed. You must provide either a URL or a path or content bytes.
- `cl.user_session.set()` : This API call sets a value in the user session dictionary. In this case, it initializes the `message_history` in the user's session with a system content message, indicating the chat's start.

Finally, the application defines the method called whenever the user sends a new message in the chat.

```

1  .on_message
2  async def main(message: str):
3      # Fetching message history from user session
4      message_history = cl.user_session.get("message_history")
5
6      # Appending user's message to message history
7      message_history.append({"role": "user", "content": message})
8
9      # Creating an empty Chainlit response message
10     msg = cl.Message(content="")
11
12     # Retry the OpenAI API call if it fails
13     for attempt in range(max_retries):
14         try:
15             # Refresh the OpenAI security token if using Azure AD
16             if openai.api_type == "azure_ad":
17                 refresh_openai_token()
18
19             # Sending the message to OpenAI and streaming the response
20             async for stream_resp in await openai.ChatCompletion.acreate(
21                 engine=engine,
22                 model=model,
23                 messages=message_history,
24                 temperature=temperature,
25                 stream=True
26             ):
27                 if stream_resp and len(stream_resp.choices) > 0:
28                     token = stream_resp.choices[0]["delta"].get("content", "")
29                     await msg.stream_token(token)
30             break
31         # Exception handling for different types of errors during the API call (Timeout, APIError

```

Here is a detailed explanation of the function steps:

- `@cl.on_message` : The `on_message` decorator registers a callback function `main(message: str)` to be called when the user submits a new message in the chat. It is the main function responsible for handling the chat logic.
- `cl.user_session.get()` : This API call retrieves a value from the user's session data stored in the `user_session` dictionary. In this case, it fetches the `message_history` from the user's session to maintain the chat history.
- `message_history.append()` : This API call appends a new message to the `message_history` list. It is used to add the user's message and the assistant's response to the chat history.
- `cl.Message()` : This API call creates a Chainlit `Message` object. The `Message` class is designed to send, stream, edit, or remove messages in the chatbot user interface. In this sample, the `Message` object is used to stream the OpenAI response in the chat.
- `msg.stream_token()` : The `stream_token` method of the `Message` class streams a token to the response message. It is used to send the response from the OpenAI Chat API in chunks to ensure real-time streaming in the chat.
- `await openai.ChatCompletion.acreate()` : This API call sends a message to the `OpenAI Chat API` in an asynchronous mode and streams the response. It uses the provided `message_history` as context for generating the assistant's response.
- The section also includes an exception handling block that retries the OpenAI API call in case of specific errors like timeouts, API errors, connection errors, invalid requests, service unavailability, and other non-

retriable errors. You can replace this code with a general-purpose retrying library for Python like [Tenacity](#).

Below, you can read the complete code of the application.

```

# Import packages
import os
import sys
import time
import openai
import random
import logging
import chainlit as cl
from azure.identity import DefaultAzureCredential
from dotenv import load_dotenv
from dotenv import dotenv_values

# Load environment variables from .env file
if os.path.exists(".env"):
    load_dotenv(override = True)
    config = dotenv_values(".env")

# Read environment variables
temperature = float(os.environ.get("TEMPERATURE", 0.9))
api_base = os.getenv("AZURE_OPENAI_BASE")
api_key = os.getenv("AZURE_OPENAI_KEY")
api_type = os.environ.get("AZURE_OPENAI_TYPE", "azure")
api_version = os.environ.get("AZURE_OPENAI_VERSION", "2023-06-01-preview")
engine = os.getenv("AZURE_OPENAI_DEPLOYMENT")
model = os.getenv("AZURE_OPENAI_MODEL")
system_content = os.getenv("AZURE_OPENAI_SYSTEM_MESSAGE", "You are a helpful assistant.")
max_retries = int(os.getenv("MAX_RETRIES", 5))
backoff_in_seconds = float(os.getenv("BACKOFF_IN_SECONDS", 1))
token_refresh_interval = int(os.getenv("TOKEN_REFRESH_INTERVAL", 1800))

# Configure OpenAI
openai.api_type = api_type
openai.api_version = api_version
openai.api_base = api_base
openai.api_key = api_key

# Set default Azure credential
default_credential = DefaultAzureCredential(
) if openai.api_type == "azure_ad" else None

# Configure a logger
logging.basicConfig(stream = sys.stdout,
                    format = '%[(asctime)s] %(filename)s:%(lineno)d} %(levelname)s - %(message)s'
                    level = logging.INFO)
logger = logging.getLogger(__name__)

def backoff(attempt : int) -> float:
    return backoff_in_seconds * 2**attempt + random.uniform(0, 1)

# Refresh the OpenAI security token every 45 minutes
def refresh_openai_token():
    token = cl.user_session.get('openai_token')
    if token == None or token.expires_on < int(time.time()) - token_refresh_interval:
        cl.user_session.set('openai_token', default_credential.get_token(
            "https://cognitiveservices.azure.com/.default"))
        openai.api_key = cl.user_session.get('openai_token').token

@cl.on_chat_start
async def start_chat():

```

```

await cl.Avatar(
    name = "Chatbot",
    url = "https://cdn-icons-png.flaticon.com/512/8649/8649595.png"
).send()
await cl.Avatar(
    name = "Error",
    url = "https://cdn-icons-png.flaticon.com/512/8649/8649595.png"
).send()
await cl.Avatar(
    name = "User",
    url = "https://media.architecturaldigest.com/photos/5f241de2c850b2a36b415024/master/w_160
).send()
cl.user_session.set(
    "message_history",
    [{"role": "system", "content": system_content}],
)

```

@cl.on_message

```

async def main(message: str):
    message_history = cl.user_session.get("message_history")
    message_history.append({"role": "user", "content": message})

    # Create the Chainlit response message
    msg = cl.Message(content = "")

    # Retry the OpenAI API call if it fails
    for attempt in range(max_retries):
        try:
            # Refresh the OpenAI security token if using Azure AD
            if openai.api_type == "azure_ad":
                refresh_openai_token()

            # Send the message to OpenAI in an asynchronous mode and stream the response
            async for stream_resp in await openai.ChatCompletion.acreate(
                engine = engine,
                model = model,
                messages = message_history,
                temperature = temperature,
                stream = True
            ):
                if stream_resp and len(stream_resp.choices) > 0:
                    token = stream_resp.choices[0]["delta"].get("content", "")
                    await msg.stream_token(token)
                break
        except openai.error.Timeout:
            # Implement exponential backoff
            wait_time = backoff(attempt)
            logger.exception(f"OpenAI API timeout occurred. Waiting {wait_time} seconds and trying")
            time.sleep(wait_time)
        except openai.error.APIError:
            # Implement exponential backoff
            wait_time = backoff(attempt)
            logger.exception(f"OpenAI API error occurred. Waiting {wait_time} seconds and trying")
            time.sleep(wait_time)
        except openai.error.APIConnectionError:
            # Implement exponential backoff
            wait_time = backoff(attempt)
            logger.exception(f"OpenAI API connection error occurred. Check your network settings,")
            time.sleep(wait_time)
        except openai.error.InvalidRequestError:
            # Implement exponential backoff
            wait_time = backoff(attempt)
            logger.exception(f"OpenAI API invalid request error occurred. Check your request")
            time.sleep(wait_time)

```

```

121         # Implement exponential backoff
122         wait_time = backoff(attempt)
123         logger.exception(f"OpenAI API invalid request. Check the documentation for the specif
124         time.sleep(wait_time)
125     except openai.error.ServiceUnavailableError:
126         # Implement exponential backoff
127         wait_time = backoff(attempt)
128         logger.exception(f"OpenAI API service unavailable. Waiting {wait_time} seconds and tr
129         time.sleep(wait_time)
130     except Exception as e:
131         logger.exception(f"A non retrievable error occurred. {e}")
132         break
133
134     message_history.append({"role": "assistant", "content": msg.content})
    await msg.send()

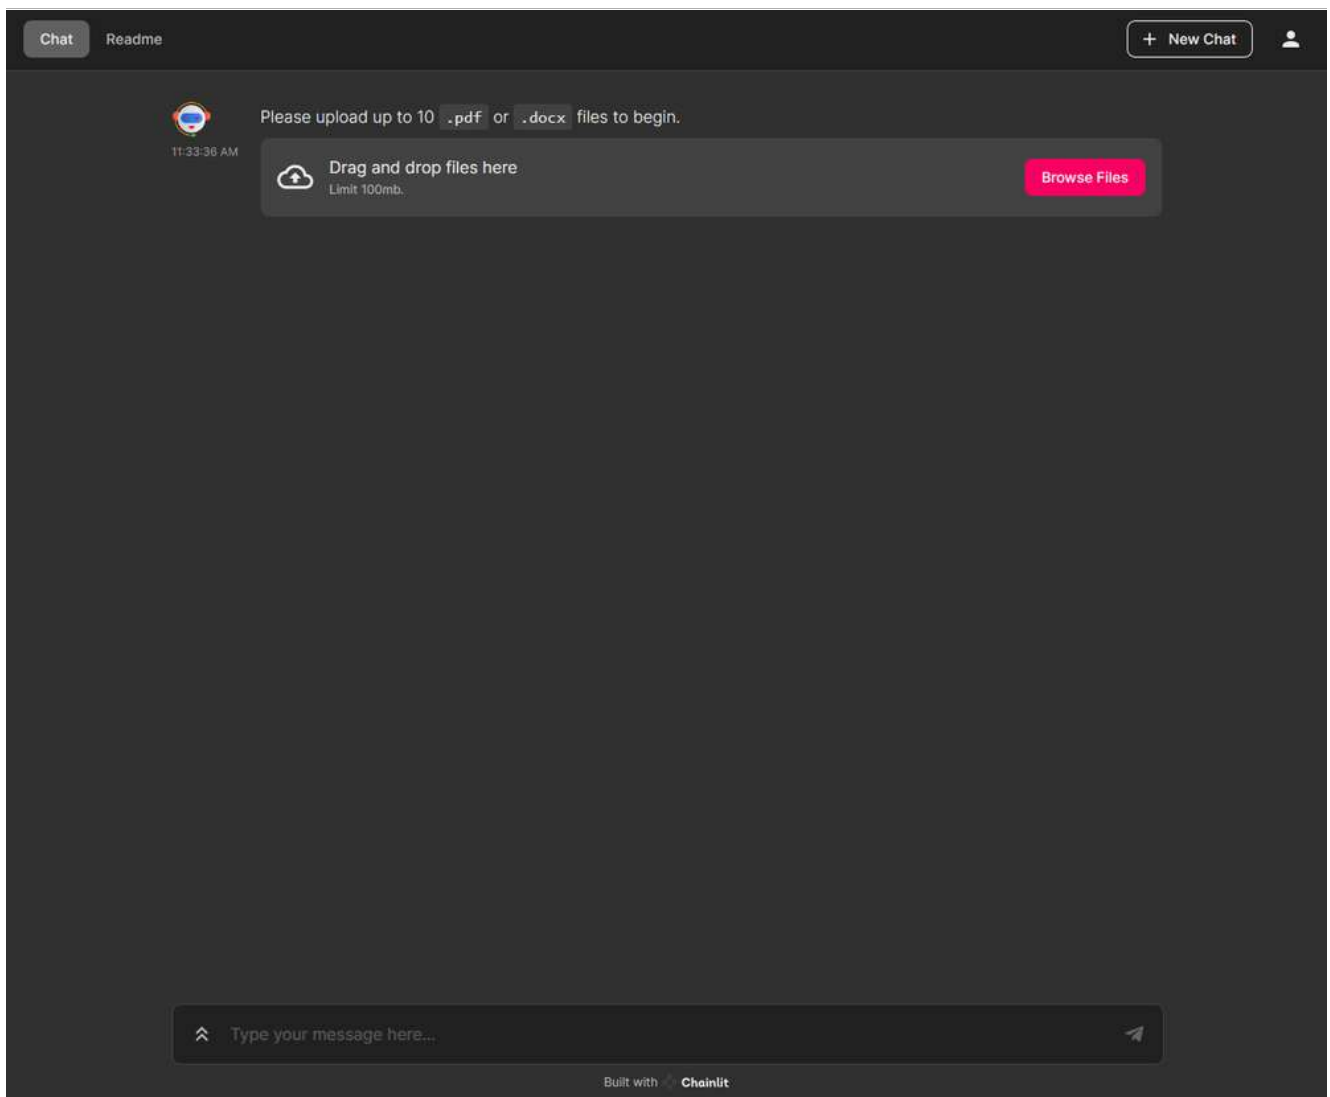
```

You can run the application locally using the following command. The `-w` flag indicates auto-reload whenever we make changes live in our application code.

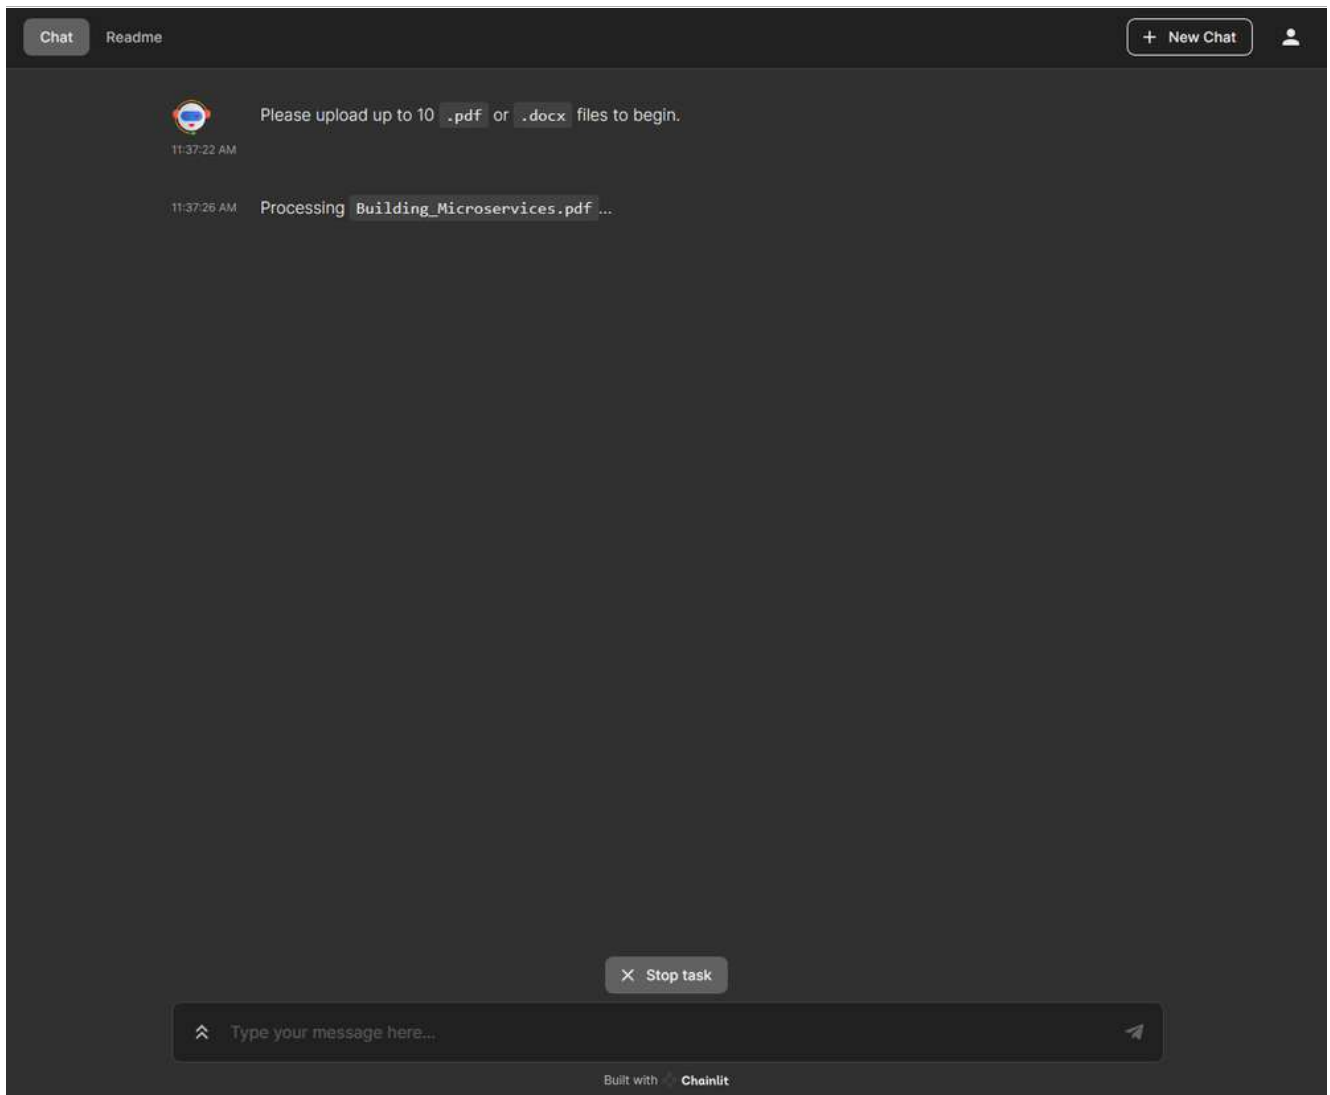
```
1 | chainlit run app.py -w
```

Documents QA Chat

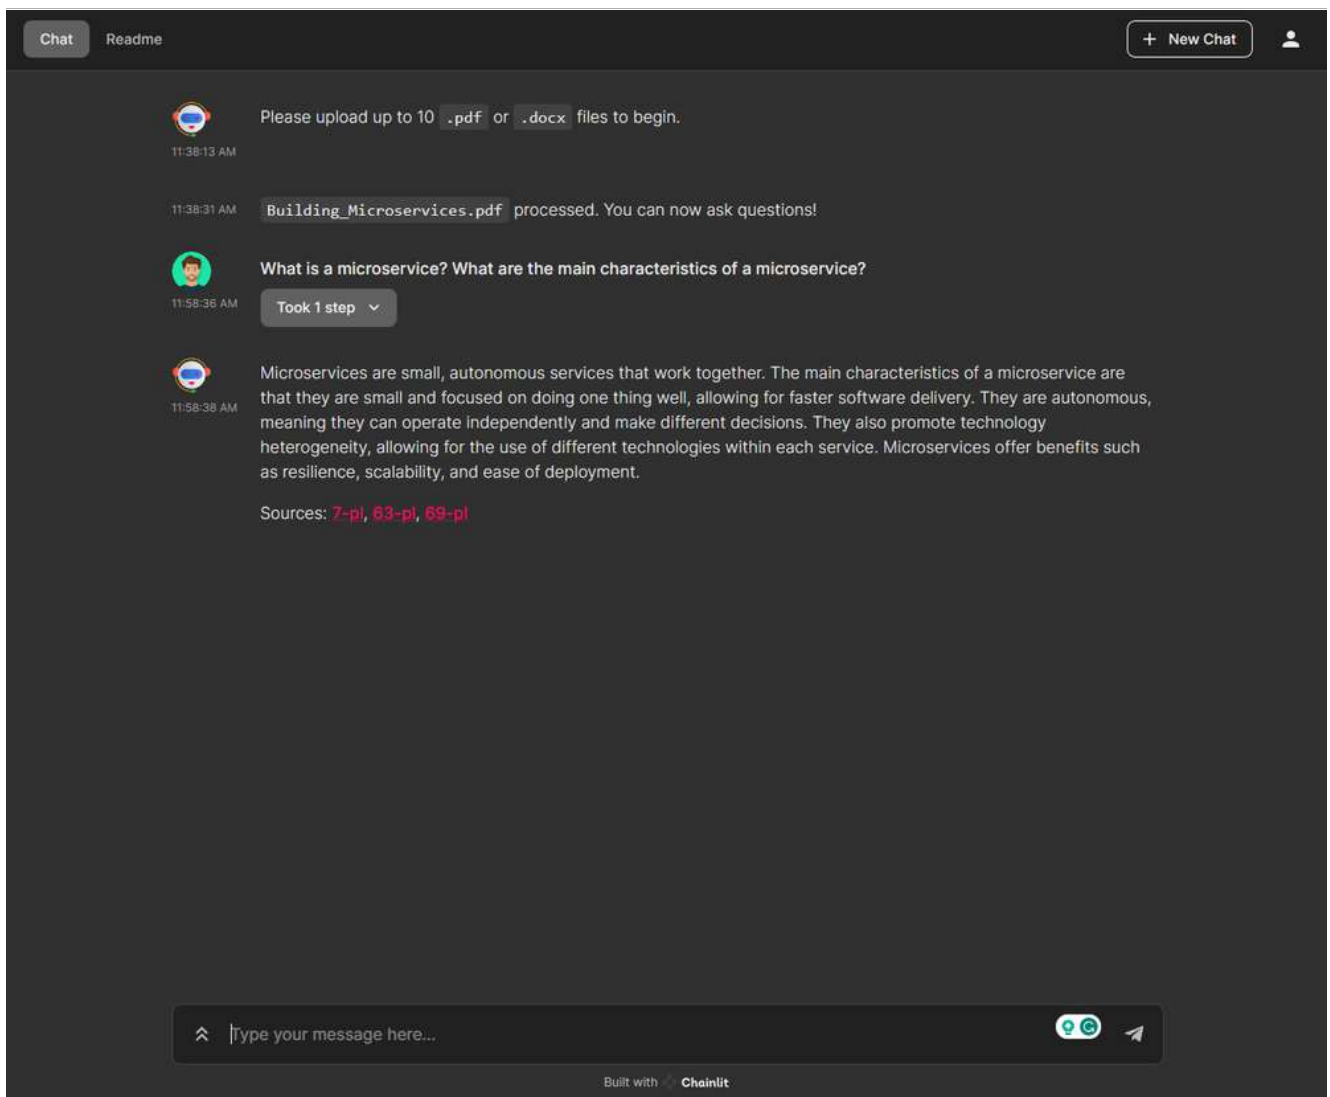
The Documents QA Chat application allows users to submit up to 10 `.pdf` and `.docx` documents. The application processes the uploaded documents to create vector embeddings. These embeddings are stored in [ChromaDB](#) vector database for efficient retrieval. Users can pose questions about the uploaded documents and view the [Chain of Thought](#), enabling easy exploration of the reasoning process. The completion message contains links to the text chunks in the files that were used as a source for the response. The following picture shows the chat application interface. As you can see, you can click the `Browse` button and choose up to 10 `.pdf` and `.docx` documents to upload. Alternatively, you can drag and drop the files over the control area.



After uploading the documents, the application creates and stores embeddings to [ChromaDB](#) vector database. During the phase, the UI shows a message Processing <file-1>, <file-2>... , as shown in the following picture:



When the code finished creating embeddings, the UI is ready to receive user's questions:



Understanding the individual steps for generating a specific answer can become challenging as your chat application grows in complexity. To solve this issue, Chainlit allows you to easily explore the reasoning process right from the user interface using the [Chain of Thought](#). If you are using the [LangChain](#) integration, every intermediary step is automatically sent and displayed in the Chainlit UI just clicking and expanding the steps, as shown in the following picture:

11:38:13 AM Please upload up to 10 .pdf or .docx files to begin.

11:38:31 AM **Building_Microservices.pdf** processed. You can now ask questions!

11:58:36 AM What is a microservice? What are the main characteristics of a microservice?

Took 1 step ^

RetrievalQ... 11:58:36 AM Took 1 step ^

StuffDocu... 11:58:36 AM Took 1 step ^

LLMChain 11:58:38 AM

Microservices are small, autonomous services that work together. The main characteristics of a microservice are that they are small and focused on doing one thing well, allowing for faster software delivery. They are autonomous, meaning they can operate independently and make different decisions. They also promote technology heterogeneity, allowing for the use of different technologies within each service. Microservices offer benefits such as resilience, scalability, and ease of deployment.
SOURCES: 7-pl, 63-pl, 69-pl

11:58:38 AM

Microservices are small, autonomous services that work together. The main characteristics of a microservice are that they are small and focused on doing one thing well, allowing for faster software delivery. They are autonomous, meaning they can operate independently and make different decisions. They also promote technology heterogeneity, allowing for the use of different technologies within each service. Microservices offer benefits such as resilience, scalability, and ease of deployment.
SOURCES: 7-pl, 63-pl, 69-pl

11:58:38 AM

Microservices are small, autonomous services that work together. The main characteristics of a microservice are that they are small and focused on doing one thing well, allowing for faster software delivery. They are autonomous, meaning they can operate independently and make different decisions. They also promote technology heterogeneity, allowing for the use of different technologies within each service. Microservices offer benefits such as resilience, scalability, and ease of deployment.
SOURCES: 7-pl, 63-pl, 69-pl

11:58:38 AM

Microservices are small, autonomous services that work together. The main characteristics of a microservice are that they are small and focused on doing one thing well, allowing for faster software delivery. They are autonomous, meaning they can operate independently and make different decisions. They also promote technology heterogeneity, allowing for the use of different technologies within each service. Microservices offer benefits such as resilience, scalability, and ease of deployment.
SOURCES: 7-pl, 63-pl, 69-pl

11:58:38 AM

Microservices are small, autonomous services that work together. The main characteristics of a microservice are that they are small and focused on doing one thing well, allowing for faster software delivery. They are autonomous, meaning they can operate independently and make different decisions. They also promote technology heterogeneity, allowing for the use of different technologies within each service. Microservices offer benefits such as resilience, scalability, and ease of deployment.
Sources: 7-pl, 63-pl, 69-pl

To see the text chunks that the large language model used to originate the response, you can click the sources links, as shown in the following picture:

Chat

Readme

+ New Chat

Please upload up to 10 .pdf or .docx files to begin.

11:38:13 AM

11:38:31 AM

Building_Microservices.pdf

processed. You can now ask questions!

What is a microservice? What are the main characteristics of a microservice?

11:58:36 AM

Took 1 step

11:58:38 AM

Microservices are small, autonomous services that work together. The main characteristics of a microservice are that they are small and focused on doing one thing well, allowing for faster software delivery. They are autonomous, meaning they can operate independently and make different decisions. They also promote technology heterogeneity, allowing for the use of different technologies within each service. Microservices offer benefits such as resilience, scalability, and ease of deployment.

Sources: [7-pl](#), [63-pl](#), [69-pl](#)

Type your message here...

Built with Chainlit

63-pl

emerged from this world. They weren't invented or described before the fact; they emerged as a trend, or a pattern, from real-world use. But they exist only because of all that has gone before. Throughout this book, I will pull strands out of this prior work to help paint a picture of how to build, manage, and evolve microservices. Many organizations have found that by embracing fine-grained, microservice architectures, they can deliver software faster and embrace newer technologies. Microservices give us significantly more freedom to react and make different decisions, allowing us to respond faster to the inevitable change that impacts all of us.

What Are Microservices?

Microservices are small, autonomous services that work together. Let's break that definition down a bit and consider the characteristics that make microservices different.

Small, and Focused on Doing One Thing Well

Codebases grow as we write code to add new features. Over time, it can be difficult to

In the Chain of Thought, below each message, you can find an `edit` button, as a pencil icon, if a prompt generated that message. Clicking on it opens the Prompt Playground dialog, allowing you to modify and iterate on the prompt as needed.

[Skip to main content](#)

Prompt playground

×

What Are Microservices? | 3Key Benefits

The benefits of microservices are many and varied. Many of these benefits can be laid out the door of any distributed system. Microservices, however, tend to achieve these benefits to a greater degree primarily due to how far they take the concepts behind distributed systems and service-oriented architecture.

Technology Heterogeneity

With a system composed of multiple, collaborating services, we can decide to use different technologies inside each one. This allows us to pick the right tool for each job, rather than having to select a more standardized, one-size-fits-all approach that often ends up being the lowest common denominator.

If one part of our system needs to improve its performance, we might decide to use a different technology stack that is better able to achieve the performance levels

Source: 69-pl

Content: We'll begin with an introduction to microservices, including the key benefits as well as some of the downsides.

xiv | PrefaceChapter 2, The Evolutionary Architect

This chapter discusses the difficulties we face in terms of making trade-offs as architects, and covers specifically just how many things we need to think about with microservices.

Chapter 3, How to Model Services

Here we'll start to define the boundary of microservices, using techniques from domain-driven design to help focus our thinking.

Chapter 4, Integration

This is where we start getting a bit deeper into specific technology implications, as we discuss what sorts of service collaboration techniques will help us most.

We'll also delve into the topic of user interfaces and integrating with legacy and commercial off-the-shelf (COTS) products.

Chapter 5, Splitting the Monolith

Many people get interested in microservices as an antidote to large, hard-to-

Source: 54-pl

What is a microservice? What are the main characteristics of a microservice?

A microservice is a small, autonomous service that works together with other microservices. The main characteristics of a microservice are that it is small and focused on doing one thing well. It is also autonomous, meaning it can independently operate and make decisions. Microservices also allow for technology heterogeneity, meaning different technologies can be used within each microservice. Additionally, microservices promote resilience and scaling, and they are easy to deploy.

SOURCES: 7-pl, 63-pl, 69-pl

Model

text-davinci-003

Temperature

0

Stop sequences

×

Top P

1

Frequency penalty

0

Presence penalty

0

Submit

↺

Let's take a look at the individual parts of the application code. The Python code starts by importing the necessary packages/modules in the following section.

```

1  # Import packages
2  import os
3  import io
4  import sys
5  import time
6  import openai
7  import random
8  import logging
9  import chainlit as cl
10 from pypdf import PdfReader
11 from docx import Document
12 from azure.identity import DefaultAzureCredential
13 from dotenv import load_dotenv
14 from dotenv import dotenv_values
15 from langchain.embeddings.openai import OpenAIEmbeddings
16 from langchain.text_splitter import RecursiveCharacterTextSplitter
17 from langchain.vectorstores import Chroma
18 from langchain.chains import RetrievalQAWithSourcesChain
19 from langchain.chat_models import AzureChatOpenAI
20 from langchain.prompts.chat import (
21     ChatPromptTemplate,
22     SystemMessagePromptTemplate,
23     HumanMessagePromptTemplate,
24 )
25
26 # These three lines swap the stdlib sqlite3 lib with the pysqlite3 package
27 __import__('pysqlite3')
28 sys.modules['sqlite3'] = sys.modules.pop('pysqlite3')
29
30 # Load environment variables from .env file
31 if os.path.exists(".env"):
32     load_dotenv(override=True)
33     config = dotenv_values(".env")

```

These are the libraries used by the chat application:

1. `os` : This module provides a way of interacting with the operating system, enabling the code to access environment variables, file paths, etc.
2. `sys` : This module provides access to some variables used or maintained by the interpreter and functions that interact with the interpreter.
3. `time` : This module provides various time-related functions for time manipulation and measurement.
4. `openai` : The OpenAI Python library provides convenient access to the OpenAI API from applications written in Python. It includes a pre-defined set of classes for API resources that initialize themselves dynamically from API responses, which makes it compatible with a wide range of versions of the OpenAI API. You can find usage examples for the OpenAI Python library in our [API reference](#) and the [OpenAI Cookbook](#).
5. `random` : This module provides functions to generate random numbers.
6. `logging` : This module provides flexible logging of messages.
7. `chainlit as cl` : This imports the [Chainlit](#) library and aliases it as `cl`. Chainlit is used to create the UI of the application.
8. `DefaultAzureCredential` from `azure.identity` : when the `openai_type` property value is `azure_ad`, a `DefaultAzureCredential` object from the [Azure Identity client library for Python - version 1.13.0](#) is

used to acquire security token from the Azure Active Directory using the credentials of the user-defined managed identity, whose client ID is defined in the `AZURE_CLIENT_ID` environment variable.

9. `load_dotenv` and `dotenv_values` from `dotenv` : [Python-dotenv](#) reads key-value pairs from a `.env` file and can set them as environment variables. It helps in the development of applications following the [12-factor](#) principles.
10. `langchain` : Large language models (LLMs) are emerging as a transformative technology, enabling developers to build applications that they previously could not. However, using these LLMs in isolation is often insufficient for creating a truly powerful app - the real power comes when you can combine them with other sources of computation or knowledge. LangChain library aims to assist in the development of those types of applications.

The `requirements.txt` file under the `src` folder contains the list of packages used by the chat applications. You can restore these packages in your environment using the following command:

```
1 | pip install -r requirements.txt --upgrade
```

Next, the code reads environment variables and configures the OpenAI settings.

```

1 # Read environment variables
2 temperature = float(os.environ.get("TEMPERATURE", 0.9))
3 api_base = os.getenv("AZURE_OPENAI_BASE")
4 api_key = os.getenv("AZURE_OPENAI_KEY")
5 api_type = os.environ.get("AZURE_OPENAI_TYPE", "azure")
6 api_version = os.environ.get("AZURE_OPENAI_VERSION", "2023-06-01-preview")
7 chat_completion_deployment = os.getenv("AZURE_OPENAI_DEPLOYMENT")
8 embeddings_deployment = os.getenv("AZURE_OPENAI_ADA_DEPLOYMENT")
9 model = os.getenv("AZURE_OPENAI_MODEL")
10 max_size_mb = int(os.getenv("CHAINLIT_MAX_SIZE_MB", 100))
11 max_files = int(os.getenv("CHAINLIT_MAX_FILES", 10))
12 text_splitter_chunk_size = int(os.getenv("TEXT_SPLITTER_CHUNK_SIZE", 1000))
13 text_splitter_chunk_overlap = int(os.getenv("TEXT_SPLITTER_CHUNK_OVERLAP", 10))
14 embeddings_chunk_size = int(os.getenv("EMBEDDINGS_CHUNK_SIZE", 16))
15 max_retries = int(os.getenv("MAX_RETRIES", 5))
16 backoff_in_seconds = float(os.getenv("BACKOFF_IN_SECONDS", 1))
17 token_refresh_interval = int(os.getenv("TOKEN_REFRESH_INTERVAL", 1800))
18
19 # Configure system prompt
20 system_template = """Use the following pieces of context to answer the users question.
21 If you don't know the answer, just say that you don't know, don't try to make up an answer.
22 ALWAYS return a "SOURCES" part in your answer.
23 The "SOURCES" part should be a reference to the source of the document from which you got your an
24
25 Example of your response should be:
26
27 ---
28
29 The answer is foo
30 SOURCES: xyz
31
32 ---
33
34 Begin!
35 -----
36 {summaries}"""
37 messages = [
38     SystemMessagePromptTemplate.from_template(system_template),
39     HumanMessagePromptTemplate.from_template("{question}"),
40 ]
41 prompt = ChatPromptTemplate.from_messages(messages)
42 chain_type_kwargs = {"prompt": prompt}
43
44 # Configure OpenAI
45 openai.api_type = api_type
46 openai.api_version = api_version
47 openai.api_base = api_base
48 openai.api_key = api_key

```

Here's a brief explanation of each variable and related environment variable:

1. `temperature` : A float value representing the temperature for Create chat completion method of the OpenAI API. It is fetched from the environment variables with a default value of 0.9.
2. `api_base` : The base URL for the OpenAI API.
3. `api_key` : The API key for the OpenAI API.
4. `api_type` : A string representing the type of the OpenAI API.
5. `api_version` : A string representing the version of the OpenAI API.

6. `chat_completion_deployment` : the name of the Azure OpenAI GPT model for chat completion.
7. `embeddings_deployment` : the name of the Azure OpenAI deployment for embeddings.
8. `model` : The model used for chat completion calls (e.g, `gpt-35-turbo-16k`).
9. `max_size_mb` : the maximum size for the uploaded documents.
10. `max_files` : the maximum number of documents that can be uploaded.
11. `text_splitter_chunk_size` : the maximum chunk size used by the `RecursiveCharacterTextSplitter` object.
12. `text_splitter_chunk_overlap` : the maximum chunk overlap used by the `RecursiveCharacterTextSplitter` object.
13. `embeddings_chunk_size` : the maximum chunk size used by the `OpenAIEmbeddings` object.
14. `max_retries` : The maximum number of retries for OpenAI API calls.
15. `backoff_in_seconds` : The backoff time in seconds for retries in case of failures.
16. `system_template` : The content of the system message used for OpenAI API calls.

In the next section, the code sets the default Azure credential based on the `api_type` and configures a logger for logging purposes.

```

1 | # Set default Azure credential
2 | default_credential = DefaultAzureCredential() if openai.api_type == "azure_ad" else None
3 |
4 | # Configure a logger
5 | logging.basicConfig(
6 |     stream=sys.stdout,
7 |     format='[%(asctime)s] %(filename)s:%(lineno)d} %(levelname)s - %(message)s',
8 |     level=logging.INFO
9 | )
10 | logger = logging.getLogger(__name__)

```

Here's a brief explanation:

1. `default_credential` : It sets the default Azure credential to `DefaultAzureCredential()` if the `api_type` is "azure_ad"; otherwise, it is set to `None` .
2. `logging.basicConfig()` : This function configures the logging system with specific settings.
 - `stream` : The output stream where log messages will be written. Here, it is set to `sys.stdout` for writing log messages to the standard output.
 - `format` : The format string for log messages. It includes the timestamp, filename, line number, log level, and the actual log message.
 - `level` : The logging level. It is set to `logging.INFO` , meaning only messages with the level `INFO` and above will be logged.
3. `logger` : This creates a logger instance named after the current module (`__name__`). The logger will be used to log messages throughout the code.

Next, the code defines a helper function `backoff` that takes an integer `attempt` and returns a float value representing the backoff time for exponential retries in case of API call failures.

```

1 | def backoff(attempt: int) -> float:
2 |     return backoff_in_seconds * 2 ** attempt + random.uniform(0, 1)

```

The backoff time is calculated using the `backoff_in_seconds` and `attempt` variables. It follows the formula `backoff_in_seconds * 2 ** attempt + random.uniform(0, 1)` . This formula increases the backoff time exponentially with each attempt and [skips to the next section](#) ~~skips to the next section~~ [skips to the next section](#) between 0 and 1 to avoid synchronized retries.

Then, the application defines a function called `refresh_openai_token()` to refresh the OpenAI security token if needed.

```
1 def refresh_openai_token():
2     token = cl.user_session.get('openai_token')
3     if token is None or token.expires_on < int(time.time()) - token_refresh_interval:
4         cl.user_session.set('openai_token', default_credential.get_token("https://cognitiveservic
5         openai.api_key = cl.user_session.get('openai_token').token
```

The function follows these steps:

1. It fetches the current token from `cl.user_session` (which seems to be a part of the `chainlit` library) using the key `'openai_token'`. The user session is a dictionary that stores the user's session data. The `id` and `env` keys are reserved for the session ID and environment variables. Other keys can be used to store arbitrary data in the user's session.
2. It checks if the token is `None` or if its expiration time (`expires_on`) is less than the current time minus 1800 seconds (30 minutes).

Next, the code defines a function called `start_chat` that is used to initialize the when the user connects to the application or clicks the `New Chat` button.

```
1 .on_chat_start
2 async def start_chat():
3     # Sending Avatars for Chat Participants
4     await cl.Avatar(
5         name="Chatbot",
6         url="https://cdn-icons-png.flaticon.com/512/8649/8649595.png"
7     ).send()
8     await cl.Avatar(
9         name="Error",
10        url="https://cdn-icons-png.flaticon.com/512/8649/8649595.png"
11    ).send()
12    await cl.Avatar(
13        name="User",
14        url="https://media.architecturaldigest.com/photos/5f241de2c850b2a36b415024/master/w_1600%
15    ).send()
```

Here is a brief explanation of the function steps:

- `@cl.on_chat_start` : The on_chat_start decorator registers a callback function `start_chat()` to be called when the Chainlit chat starts. It is used to set up the chat and send avatars for the Chatbot, Error, and User participants in the chat.
- `cl.Avatar()` : the Avatar class allows you to display an avatar image next to a message instead of the author's name. You need to send the element once. Next, if an avatar's name matches an author's name, the avatar will be automatically displayed. You must provide either a URL or a path or content bytes.

The following code is used to initialize the large language model (LLM) chain used to reply to questions on the content of the uploaded documents.

```

1 | # Initialize the file list to None
2 |     files = None
3 |
4 |     # Wait for the user to upload a file
5 |     while files is None:
6 |         files = await cl.AskFileMessage(
7 |             content=f"Please upload up to {max_files} `.pdf` or `.docx` files to begin.",
8 |             accept=["application/pdf", "application/vnd.openxmlformats-officedocument.wordprocess
9 |                 max_size_mb=max_size_mb,
10 |                 max_files=max_files,
11 |                 timeout=86400,
12 |                 raise_on_timeout=False
13 |             ).send()

```

The [AskFileMessage](#) API call prompts the user to upload up to a specified number of .pdf or .docx files. The uploaded files are stored in the `files` variable. The process continues until the user uploads files. For more information, see [AskFileMessage](#).

The following code processes each uploaded file by extracting its content.

1. The text content of each file is stored in the list `all_texts`.
2. This code performs text processing and chunking. It checks the file extension to read the file content accordingly, depending on if it's a .pdf or a .docx document.
3. The text content is split into smaller chunks using the [RecursiveCharacterTextSplitter](#) LangChain object.
4. Metadata is created for each chunk and stored in the `metadatas` list.
5. If `openai.api_type == "azure_ad"`, the code invokes the `refresh_openai_token()` that gets a security token from Azure AD to communicate with the Azure OpenAI Service.


```

1  # Create a message to inform the user that the files are being processed
2  content = ''
3  if (len(files) == 1):
4      content = f"Processing `{files[0].name}`..."
5  else:
6      files_names = [f"`{f.name}`" for f in files]
7      content = f"Processing {' '.join(files_names)}..."
8  msg = cl.Message(content = content, author = "Chatbot")
9  await msg.send()
10
11 # Create a list to store the texts of each file
12 all_texts = []
13
14 # Process each file uploaded by the user
15 for file in files:
16
17     # Create an in-memory buffer from the file content
18     bytes = io.BytesIO(file.content)
19
20     # Get file extension
21     extension = file.name.split('.')[1]
22
23     # Initialize the text variable
24     text = ''
25
26     # Read the file
27     if extension == "pdf":
28         # ...
29     elif extension == "docx":
30         # ...
31
32     # Split the text into chunks
33     text_splitter = RecursiveCharacterTextSplitter(
34         chunk_size=text_splitter_chunk_size,
35         chunk_overlap=text_splitter_chunk_overlap
36     )
37     texts = text_splitter.split_text(text)
38
39     # Add the chunks and metadata to the list
40     all_texts.extend(texts)
41
42     # Create a metadata for each chunk
43     metadatas = [{"source": f"{i}-pl"} for i in range(len(all_texts))]
44
45     # Refresh the OpenAI security token if using Azure AD
46     if openai.api_type == "azure_ad":
47         refresh_openai_token()

```

The next piece of code performs the following steps:

1. It creates an [OpenAIEmbeddings](#) configured to use the embeddings model in the Azure OpenAI Service to create embeddings from text chunks.
2. It creates a [ChromaDB](#) vector database using the `OpenAIEmbeddings` object, the text chunks list, and the metadata list.
3. It creates an [AzureChatOpenAI](#) LangChain object based on the GPT model hosted in Azure OpenAI Service.

4. It creates a chain using the `RetrievalQAWithSourcesChain.from_chain_type` API call uses previously created models and stores them as retrievers.
5. It stores the metadata and text chunks in the user session using the `c1.user_session.set()` API call.
6. It creates a message to inform the user that the files are ready for queries, and finally returns the `chain`.
7. The `c1.user_session.set("chain", chain)` call stores the LLM chain in the `user_session` dictionary for later use.

```

1  # Refresh the OpenAI security token if using Azure AD
2  if openai.api_type == "azure_ad":
3      refresh_openai_token()
4
5  # Create a Chroma vector store
6  embeddings = OpenAIEmbeddings(
7      deployment = embeddings_deployment,
8      openai_api_key = openai.api_key,
9      openai_api_base = openai.api_base,
10     openai_api_version = openai.api_version,
11     openai_api_type = openai.api_type,
12     chunk_size = embeddings_chunk_size)
13
14  # Create a Chroma vector store
15  db = await cl.make_async(Chroma.from_texts)(
16     all_texts, embeddings, metadatas = metadatas
17  )
18
19  # Create an AzureChatOpenAI llm
20  llm = AzureChatOpenAI(
21     temperature = temperature,
22     openai_api_key = openai.api_key,
23     openai_api_base = openai.api_base,
24     openai_api_version = openai.api_version,
25     openai_api_type = openai.api_type,
26     deployment_name = chat_completion_deployment)
27
28  # Create a chain that uses the Chroma vector store
29  chain = RetrievalQAWithSourcesChain.from_chain_type(
30     llm = llm,
31     chain_type = "stuff",
32     retriever = db.as_retriever(),
33     return_source_documents = True,
34     chain_type_kwargs = chain_type_kwargs
35  )
36
37  # Save the metadata and texts in the user session
38  cl.user_session.set("metadatas", metadatas)
39  cl.user_session.set("texts", all_texts)
40
41  # Create a message to inform the user that the files are ready for queries
42  content = ''
43  if (len(files) == 1):
44     content = f"`{files[0].name}` processed. You can now ask questions!"
45  else:
46     files_names = [f"`{f.name}`" for f in files]
47     content = f'{' , '.join(files_names)} processed. You can now ask questions.'
48  msg.content = content
49  msg.author = "Chatbot"
50  await msg.update()
51
52  # Store the chain in the user session
53  cl.user_session.set("chain", chain)

```

The following code handles the communication with the OpenAI API and incorporates retrying logic in case the API calls fail due to specific errors.

- `@cl.on_message` : The `on_message` decorator registers a callback function `main(message: str)` to be called when the user submits a new `Message` in the chatbot. The main function responsible for handling

the chat logic.

- `cl.user_session.get("chain")` : this call retrieves the LLM chain from the user_session dictionary.
- The `for` loop allows multiple attempts, up to `max_retries` , to communicate with the chat completion API and handles different types of API errors, such as timeout, connection error, invalid request, and service unavailability.
- `await chain.acall` : The asynchronous call to the RetrievalQAWithSourcesChain.acall executes the LLM chain with the user message as an input.

```

1  .on_message
2  async def run(message: str):
3      # Retrieve the chain from the user session
4      chain = cl.user_session.get("chain")
5
6      # Initialize the response
7      response = None
8
9      # Retry the OpenAI API call if it fails
10     for attempt in range(max_retries):
11         try:
12             # Refresh the OpenAI security token if using Azure AD
13             if openai.api_type == "azure_ad":
14                 refresh_openai_token()
15
16             # Ask the question to the chain
17             response = await chain.acall(message, callbacks=[cl.AsyncLangchainCallbackHandler()])
18             break
19         except openai.error.Timeout:
20             # Exception handling for timeout error
21             # Implement exponential backoff
22             wait_time = backoff(attempt)
23             logger.exception(f"OpenAI API timeout occurred. Waiting {wait_time} seconds and trying")
24             time.sleep(wait_time)
25         except openai.error.APIError:
26             # Exception handling for API error
27             # Implement exponential backoff
28             wait_time = backoff(attempt)
29             logger.exception(f"OpenAI API error occurred. Waiting {wait_time} seconds and trying")
30             time.sleep(wait_time)
31         except openai.error.APIConnectionError:
32             # Exception handling for API connection error
33             # Implement exponential backoff
34             wait_time = backoff(attempt)
35             logger.exception(f"OpenAI API connection error occurred. Check your network settings,")
36             time.sleep(wait_time)
37         except openai.error.InvalidRequestError:
38             # Exception handling for invalid request error
39             # Implement exponential backoff
40             wait_time = backoff(attempt)
41             logger.exception(f"OpenAI API invalid request. Check the documentation for the specif")
42             time.sleep(wait_time)
43         except openai.error.ServiceUnavailableError:
44             # Exception handling for service unavailable error
45             # Implement exponential backoff
46             wait_time = backoff(attempt)
47             logger.exception(f"OpenAI API service unavailable. Waiting {wait_time} seconds and tr")
48             time.sleep(wait_time)
49         except Exception as e:
50             # Exception handling for non-retriable errors
51             logger.exception(f"A non-retriable error occurred. {e}")
52             break

```

The code below extracts the answers and sources from the API response and formats them to be sent as a message.

- The answer and sources are obtained from the response dictionary.

- The sources are then processed to find corresponding texts in the user session metadata (`metadatas`) and create `source_elements` using `cl.Text()` .
- `cl.Message().send()` : the Message API creates and displays a message containing the answer and sources, if available.

```

1  # Get the answer and sources from the response
2      answer = response["answer"]
3      sources = response["sources"].strip()
4      source_elements = []
5
6      # Get the metadata and texts from the user session
7      metadatas = cl.user_session.get("metadatas")
8      all_sources = [m["source"] for m in metadatas]
9      texts = cl.user_session.get("texts")
10
11     if sources:
12         found_sources = []
13
14         # Add the sources to the message
15         for source in sources.split(","):
16             source_name = source.strip().replace(".", "")
17             # Get the index of the source
18             try:
19                 index = all_sources.index(source_name)
20             except ValueError:
21                 continue
22             text = texts[index]
23             found_sources.append(source_name)
24             # Create the text element referenced in the message
25             source_elements.append(cl.Text(content=text, name=source_name))
26
27     if found_sources:
28         answer += f"\nSources: {'', '.join(found_sources)}"
29     else:
30         answer += "\nNo sources found"
31
32     await cl.Message(content=answer, elements=source_elements).send()

```

Below, you can read the complete code of the application.

```

# Import packages
import os
import io
import sys
import time
import openai
import random
import logging
import chainlit as cl
from pypdf import PdfReader
from docx import Document
from azure.identity import DefaultAzureCredential
from dotenv import load_dotenv
from dotenv import dotenv_values
from langchain.embeddings.openai import OpenAIEmbeddings
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain.vectorstores import Chroma
from langchain.chains import RetrievalQAWithSourcesChain
from langchain.chat_models import AzureChatOpenAI
from langchain.prompts.chat import (
    ChatPromptTemplate,
    SystemMessagePromptTemplate,
    HumanMessagePromptTemplate,
)

# These three lines swap the stdlib sqlite3 lib with the pysqlite3 package
__import__('pysqlite3')
sys.modules['sqlite3'] = sys.modules.pop('pysqlite3')

# Load environment variables from .env file
if os.path.exists(".env"):
    load_dotenv(override = True)
    config = dotenv_values(".env")

# Read environment variables
temperature = float(os.environ.get("TEMPERATURE", 0.9))
api_base = os.getenv("AZURE_OPENAI_BASE")
api_key = os.getenv("AZURE_OPENAI_KEY")
api_type = os.environ.get("AZURE_OPENAI_TYPE", "azure")
api_version = os.environ.get("AZURE_OPENAI_VERSION", "2023-06-01-preview")
chat_completion_deployment = os.getenv("AZURE_OPENAI_DEPLOYMENT")
embeddings_deployment = os.getenv("AZURE_OPENAI_ADA_DEPLOYMENT")
model = os.getenv("AZURE_OPENAI_MODEL")
max_size_mb = int(os.getenv("CHAINLIT_MAX_SIZE_MB", 100))
max_files = int(os.getenv("CHAINLIT_MAX_FILES", 10))
text_splitter_chunk_size = int(os.getenv("TEXT_SPLITTER_CHUNK_SIZE", 1000))
text_splitter_chunk_overlap = int(os.getenv("TEXT_SPLITTER_CHUNK_OVERLAP", 10))
embeddings_chunk_size = int(os.getenv("EMBEDDINGS_CHUNK_SIZE", 16))
max_retries = int(os.getenv("MAX_RETRIES", 5))
backoff_in_seconds = float(os.getenv("BACKOFF_IN_SECONDS", 1))

# Configure system prompt
system_template = """Use the following pieces of context to answer the users question.
If you don't know the answer, just say that you don't know, don't try to make up an answer.
ALWAYS return a "SOURCES" part in your answer.
The "SOURCES" part should be a reference to the source of the document from which you got your an

Example of your response should be:

```

The answer is foo

SOURCES: xyz

Begin!

```
{summaries}"""
```

```
messages = [
```

```
    SystemMessagePromptTemplate.from_template(system_template),
```

```
    HumanMessagePromptTemplate.from_template("{question}"),
```

```
]
```

```
prompt = ChatPromptTemplate.from_messages(messages)
```

```
chain_type_kwargs = {"prompt": prompt}
```

```
# Configure OpenAI
```

```
openai.api_type = api_type
```

```
openai.api_version = api_version
```

```
openai.api_base = api_base
```

```
openai.api_key = api_key
```

```
# Set default Azure credential
```

```
default_credential = DefaultAzureCredential(
```

```
) if openai.api_type == "azure_ad" else None
```

```
# Configure a logger
```

```
logging.basicConfig(stream = sys.stdout,
```

```
                    format = '%(asctime)s] %(filename)s:%(lineno)d} %(levelname)s - %(message)s
```

```
                    level = logging.INFO)
```

```
logger = logging.getLogger(__name__)
```

```
# Refresh the OpenAI security token every 45 minutes
```

```
def refresh_openai_token():
```

```
    token = cl.user_session.get('openai_token')
```

```
    if token == None or token.expires_on < int(time.time()) - 1800:
```

```
        cl.user_session.set('openai_token', default_credential.get_token(
            "https://cognitiveservices.azure.com/.default"))
```

```
        openai.api_key = cl.user_session.get('openai_token').token
```

```
def backoff(attempt : int) -> float:
```

```
    return backoff_in_seconds * 2**attempt + random.uniform(0, 1)
```

```
@cl.on_chat_start
```

```
async def start():
```

```
    await cl.Avatar(
```

```
        name = "Chatbot",
```

```
        url = "https://cdn-icons-png.flaticon.com/512/8649/8649595.png"
```

```
    ).send()
```

```
    await cl.Avatar(
```

```
        name = "Error",
```

```
        url = "https://cdn-icons-png.flaticon.com/512/8649/8649595.png"
```

```
    ).send()
```

```
    await cl.Avatar(
```

```
        name = "User",
```

```
        url = "https://media.architecturaldigest.com/photos/5f241de2c850b2a36b415024/master/w_160
```

```
    ).send()
```

```
# Initialize the file list to skip for later navigation
```



```

files = None

# Wait for the user to upload a file
while files == None:
    files = await cl.AskFileMessage(
        content = f"Please upload up to {max_files} `.pdf` or `.docx` files to begin.",
        accept = ["application/pdf",
            "application/vnd.openxmlformats-officedocument.wordprocessingml.document"],
        max_size_mb = max_size_mb,
        max_files = max_files,
        timeout = 86400,
        raise_on_timeout = False
    ).send()

# Create a message to inform the user that the files are being processed
content = ''
if (len(files) == 1):
    content = f"Processing `{files[0].name}`..."
else:
    files_names = [f"`{f.name}`" for f in files]
    content = f"Processing {', '.join(files_names)}..."
msg = cl.Message(content = content, author = "Chatbot")
await msg.send()

# Create a list to store the texts of each file
all_texts = []

# Process each file uploaded by the user
for file in files:

    # Create an in-memory buffer from the file content
    bytes = io.BytesIO(file.content)

    # Get file extension
    extension = file.name.split('.')[-1]

    # Initialize the text variable
    text = ''

    # Read the file
    if extension == "pdf":
        reader = PdfReader(bytes)
        for i in range(len(reader.pages)):
            text += reader.pages[i].extract_text()
    elif extension == "docx":
        doc = Document(bytes)
        paragraph_list = []
        for paragraph in doc.paragraphs:
            paragraph_list.append(paragraph.text)
        text = '\n'.join(paragraph_list)

    # Split the text into chunks
    text_splitter = RecursiveCharacterTextSplitter(
        chunk_size = text_splitter_chunk_size,
        chunk_overlap = text_splitter_chunk_overlap)
    texts = text_splitter.split_text(text)

    # Add the chunks and metadata to the list
    all_texts.extend(texts)

```

```
# Create a metadata for each chunk
metadatas = [{"source": f"{i}-pl"} for i in range(len(all_texts))]

# Refresh the OpenAI security token if using Azure AD
if openai.api_type == "azure_ad":
    refresh_openai_token()

# Create a Chroma vector store
embeddings = OpenAIEmbeddings(
    deployment = embeddings_deployment,
    openai_api_key = openai.api_key,
    openai_api_base = openai.api_base,
    openai_api_version = openai.api_version,
    openai_api_type = openai.api_type,
    chunk_size = embeddings_chunk_size)

# Create a Chroma vector store
db = await cl.make_async(Chroma.from_texts)(
    all_texts, embeddings, metadatas = metadatas
)

# Create an AzureChatOpenAI llm
llm = AzureChatOpenAI(
    temperature = temperature,
    openai_api_key = openai.api_key,
    openai_api_base = openai.api_base,
    openai_api_version = openai.api_version,
    openai_api_type = openai.api_type,
    deployment_name = chat_completion_deployment)

# Create a chain that uses the Chroma vector store
chain = RetrievalQAWithSourcesChain.from_chain_type(
    llm = llm,
    chain_type = "stuff",
    retriever = db.as_retriever(),
    return_source_documents = True,
    chain_type_kwargs = chain_type_kwargs
)

# Save the metadata and texts in the user session
cl.user_session.set("metadatas", metadatas)
cl.user_session.set("texts", all_texts)

# Create a message to inform the user that the files are ready for queries
content = ''
if (len(files) == 1):
    content = f"`{files[0].name}` processed. You can now ask questions!"
else:
    files_names = [f"`{f.name}`" for f in files]
    content = f'{{', '.join(files_names)}} processed. You can now ask questions.'
msg.content = content
msg.author = "Chatbot"
await msg.update()

# Store the chain in the user session
cl.user_session.set("chain", chain)

@cl.on_message
async def run(message: str):
    # Retrieve the chain from the user session for re-creation
```

```

chain = cl.user_session.get("chain")

# Initialize the response
response = None

# Retry the OpenAI API call if it fails
for attempt in range(max_retries):
    try:
        # Refresh the OpenAI security token if using Azure AD
        if openai.api_type == "azure_ad":
            refresh_openai_token()

        # Ask the question to the chain
        response = await chain.acall(message, callbacks = [cl.AsyncLangchainCallbackHandler()])
        break
    except openai.error.Timeout:
        # Implement exponential backoff
        wait_time = backoff(attempt)
        logger.exception(f"OpenAI API timeout occurred. Waiting {wait_time} seconds and trying again")
        time.sleep(wait_time)
    except openai.error.APIError:
        # Implement exponential backoff
        wait_time = backoff(attempt)
        logger.exception(f"OpenAI API error occurred. Waiting {wait_time} seconds and trying again")
        time.sleep(wait_time)
    except openai.error.APIConnectionError:
        # Implement exponential backoff
        wait_time = backoff(attempt)
        logger.exception(f"OpenAI API connection error occurred. Check your network settings, and try again")
        time.sleep(wait_time)
    except openai.error.InvalidRequestError:
        # Implement exponential backoff
        wait_time = backoff(attempt)
        logger.exception(f"OpenAI API invalid request. Check the documentation for the specific error")
        time.sleep(wait_time)
    except openai.error.ServiceUnavailableError:
        # Implement exponential backoff
        wait_time = backoff(attempt)
        logger.exception(f"OpenAI API service unavailable. Waiting {wait_time} seconds and trying again")
        time.sleep(wait_time)
    except Exception as e:
        logger.exception(f"A non retrievable error occurred. {e}")
        break

# Get the answer and sources from the response
answer = response["answer"]
sources = response["sources"].strip()
source_elements = []

# Get the metadata and texts from the user session
metadatas = cl.user_session.get("metadatas")
all_sources = [m["source"] for m in metadatas]
texts = cl.user_session.get("texts")

if sources:
    found_sources = []

    # Add the sources to the message
    for source in sources.split(","):
        source_name = source.strip()
        source_text = cl.user_session.get("texts").get(source_name)

```

```

303         # Get the index of the source
304         try:
305             index = all_sources.index(source_name)
306         except ValueError:
307             continue
308         text = texts[index]
309         found_sources.append(source_name)
310         # Create the text element referenced in the message
311         source_elements.append(cl.Text(content = text, name = source_name))
312
313     if found_sources:
314         answer += f"\nSources: {' '.join(found_sources)}"
315     else:
316         answer += "\nNo sources found"
317
318     await cl.Message(content = answer, elements = source_elements).send()

```

You can run the application locally using the following command. The `-w` flag indicates auto-reload whenever we make changes live in our application code.

```
1 | chainlit run app.py -w
```

Build Docker Images

You can use the `src/01-build-docker-images.sh` Bash script to build the Docker container image for each container app.

```

1  #!/bin/bash
2
3  # Variables
4  source ./00-variables.sh
5
6  # Use a for loop to build the docker images using the array index
7  for index in ${!images[@]}; do
8      # Build the docker image
9      docker build -t ${images[$index]}:$tag -f Dockerfile --build-arg FILENAME=${filenames[$index]}
10 done

```

Before running any script in the `src` folder, make sure to customize the value of the variables inside the `00-variables.sh` file located in the same folder. This file is embedded in all the scripts and contains the following variables:

```

1  # Variables
2
3  # Azure Container Registry
4  prefix="Blue"
5  acrName="${prefix}Registry"
6  acrResourceGrougName="${prefix}RG"
7  location="EastUS"
8
9  # Python Files
10 docAppFile="doc.py"
11 chatAppFile="chat.py"
12
13 # Docker Images
14 docImageName="doc"
15 chatImageName="chat"
16 tag="v1"
17 port="8000"
18
19 # Arrays
20 images=($docImageName $chatImageName)
21 filenames=($docAppFile $chatAppFile)

```

The Dockerfile under the src folder is parametric and can be used to build the container images for both chat applications.

```

# app/Dockerfile

# # Stage 1 - Install build dependencies

# A Dockerfile must start with a FROM instruction that sets the base image for the container.
# The Python images come in many flavors, each designed for a specific use case.
# The python:3.11-slim image is a good base image for most applications.
# It is a minimal image built on top of Debian Linux and includes only the necessary packages to
# The slim image is a good choice because it is small and contains only the packages needed to ru
# For more information, see:
# * https://hub.docker.com/\_/python
# * https://docs.streamlit.io/knowledge-base/tutorials/deploy/docker
FROM python:3.11-slim AS builder

# The WORKDIR instruction sets the working directory for any RUN, CMD, ENTRYPOINT, COPY and ADD i
# If the WORKDIR doesn't exist, it will be created even if it's not used in any subsequent Docker
# For more information, see: https://docs.docker.com/engine/reference/builder/#workdir
WORKDIR /app

# Set environment variables.
# The ENV instruction sets the environment variable <key> to the value <value>.
# This value will be in the environment of all "descendant" Dockerfile commands and can be replac
# For more information, see: https://docs.docker.com/engine/reference/builder/#env
ENV PYTHONDONTWRITEBYTECODE 1
ENV PYTHONUNBUFFERED 1

# Install git so that we can clone the app code from a remote repo using the RUN instruction.
# The RUN comand has 2 forms:
# * RUN <command> (shell form, the command is run in a shell, which by default is /bin/sh -c on L
# * RUN ["executable", "param1", "param2"] (exec form)
# The RUN instruction will execute any commands in a new layer on top of the current image and co
# The resulting committed image will be used for the next step in the Dockerfile.
# For more information, see: https://docs.docker.com/engine/reference/builder/#run
RUN apt-get update && apt-get install -y \
    build-essential \
    curl \
    software-properties-common \
    git \
    && rm -rf /var/lib/apt/lists/*

# Create a virtualenv to keep dependencies together
RUN python -m venv /opt/venv
ENV PATH="/opt/venv/bin:$PATH"

# Clone the requirements.txt which contains dependencies to WORKDIR
# COPY has two forms:
# * COPY <src> <dest> (this copies the files from the local machine to the container's own filesy
# * COPY ["<src>",... "<dest>"] (this form is required for paths containing whitespace)
# For more information, see: https://docs.docker.com/engine/reference/builder/#copy
COPY requirements.txt .

# Install the Python dependencies
RUN pip install --no-cache-dir --no-deps -r requirements.txt

# Stage 2 - Copy only necessary files to the runner stage

# The FROM instruction initializes a new build stage for the application
FROM python:3.11-slim

```

```

61 # Define the filename to copy as an argument
62 ARG FILENAME
63
64 # Deefine the port to run the application on as an argument
65 ARG PORT=8000
66
67 # Set an environment variable
68 ENV FILENAME=${FILENAME}
69
70 # Sets the working directory to /app
71 WORKDIR /app
72
73 # Copy the virtual environment from the builder stage
74 COPY --from=builder /opt/venv /opt/venv
75
76 # Set environment variables
77 ENV PATH="/opt/venv/bin:$PATH"
78
79 # Clone the $FILENAME containing the application code
80 COPY $FILENAME .
81
82 # Copy the chainlit.md file to the working directory
83 COPY chainlit.md .
84
85 # Copy the .chainlit folder to the working directory
86 COPY ./chainlit ./chainlit
87
88 # The EXPOSE instruction informs Docker that the container listens on the specified network ports
89 # For more information, see: https://docs.docker.com/engine/reference/builder/#expose
90 EXPOSE $PORT
91
92 # The ENTRYPOINT instruction has two forms:
93 # * ENTRYPOINT ["executable", "param1", "param2"] (exec form, preferred)
94 # * ENTRYPOINT command param1 param2 (shell form)
95 # The ENTRYPOINT instruction allows you to configure a container that will run as an executable.
96 # For more information, see: https://docs.docker.com/engine/reference/builder/#entrypoint
  CMD chainlit run $FILENAME --port=$PORT

```

Test applications locally

You can use the `src/02-run-docker-container.sh` Bash script to test the containers for the sender, processor, and receiver applications.

```

1  #!/bin/bash
2
3  # Variables
4  source ./00-variables.sh
5
6  # Print the menu
7  echo "===== "
8  echo "Run Docker Container (1-3): "
9  echo "===== "
10 options=(
11     "Doc"
12     "Chat"
13 )
14 name=""
15 # Select an option
16 COLUMNS=0
17 select option in "${options[@]}"; do
18     case $option in
19         "Doc")
20             docker run -it \
21                 --rm \
22                 -p $port:$port \
23                 -e AZURE_OPENAI_BASE=$AZURE_OPENAI_BASE \
24                 -e AZURE_OPENAI_KEY=$AZURE_OPENAI_KEY \
25                 -e AZURE_OPENAI_MODEL=$AZURE_OPENAI_MODEL \
26                 -e AZURE_OPENAI_DEPLOYMENT=$AZURE_OPENAI_DEPLOYMENT \
27                 -e AZURE_OPENAI_ADA_DEPLOYMENT=$AZURE_OPENAI_ADA_DEPLOYMENT \
28                 -e AZURE_OPENAI_VERSION=$AZURE_OPENAI_VERSION \
29                 -e AZURE_OPENAI_TYPE=$AZURE_OPENAI_TYPE \
30                 -e TEMPERATURE=$TEMPERATURE \
31                 --name $docImageName \
32                 $docImageName:$tag
33             break
34         ;;
35         "Chat")
36             docker run -it \
37                 --rm \
38                 -p $port:$port \
39                 -e AZURE_OPENAI_BASE=$AZURE_OPENAI_BASE \
40                 -e AZURE_OPENAI_KEY=$AZURE_OPENAI_KEY \
41                 -e AZURE_OPENAI_MODEL=$AZURE_OPENAI_MODEL \
42                 -e AZURE_OPENAI_DEPLOYMENT=$AZURE_OPENAI_DEPLOYMENT \
43                 -e AZURE_OPENAI_VERSION=$AZURE_OPENAI_VERSION \
44                 -e AZURE_OPENAI_TYPE=$AZURE_OPENAI_TYPE \
45                 -e TEMPERATURE=$TEMPERATURE \
46                 --name $chatImageName \
47                 $chatImageName:$tag
48             break
49         ;;
50         "Quit")
51             exit
52         ;;
53         *) echo "invalid option $REPLY" ;;
54     esac
55 done

```


Push Docker containers to the Azure Container Registry

You can use the `src/03-push-docker-image.sh` Bash script to push the Docker container images for the `sender`, `processor`, and `receiver` applications to the [Azure Container Registry \(ACR\)](#).

```
1  #!/bin/bash
2
3  # Variables
4  source ./00-variables.sh
5
6  # Login to ACR
7  echo "Logging in to [{acrName,,}] container registry..."
8  az acr login --name ${acrName,,}
9
10 # Retrieve ACR login server. Each container image needs to be tagged with the loginServer name of
11 echo "Retrieving login server for the [{acrName,,}] container registry..."
12 loginServer=$(az acr show --name ${acrName,,} --query loginServer --output tsv)
13
14 # Use a for loop to tag and push the local docker images to the Azure Container Registry
15 for index in ${!images[@]}; do
16     # Tag the local sender image with the loginServer of ACR
17     docker tag ${images[$index],,}:$tag $loginServer/${images[$index],,}:$tag
18
19     # Push the container image to ACR
20     docker push $loginServer/${images[$index],,}:$tag
21 done
```

Monitoring

[Azure Container Apps](#) provides several built-in observability features that together give you a holistic view of your container app's health throughout its application lifecycle. These features help you monitor and diagnose the state of your app to improve performance and respond to trends and critical problems.

You can use the `Log Stream` panel on the Azure Portal to see the logs generated by a container app, as shown in the following screenshot.

Microsoft Azure (Preview) Search resources, services, and docs (G+)

Home > chatapp

chatapp | Log stream

Container App

Search Refresh

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Application

Revisions

Containers

Scale and replicas

Settings

Authentication

Secrets

Ingress

Continuous deployment

Custom domains

Dapr

Identity

Service Connector (preview)

CORS (preview)

Locks

Monitoring

Alerts

Metrics

Logs

Log stream

Console

Advisor recommendations

Help

Logs Console System

Replica chatapp--39jy3n9-55db787494-zmt2d

Container chat

Stop Copy Clear Maximize

```

g to the container 'chat'...
2023-07-27T09:43:30.77227 Successfully Connected to container: 'chat' [Revision: 'chatapp--39jy3n9', Replica: 'chatapp--39jy3n9-55db787494-zmt2d']
2023-07-26T23:35:16.1422823742 2023-07-26 23:35:16 - Incomplete environment configuration. These variables are set: AZURE_CLIENT_ID
2023-07-26T23:35:16.1423294302 2023-07-26 23:35:16 - ManagedIdentityCredential will use App Service managed identity
2023-07-26T23:35:16.2242655922 2023-07-26 23:35:16 - Your app is available at http://localhost:8000
2023-07-27T09:43:42.4219744722 2023-07-27 09:43:42 - Request URL: 'http://localhost:42356/msi/token?api-version=REDACTED&resource=REDACTED&client_id=REDACTED'
2023-07-27T09:43:42.4223296022 Request method: 'GET'
2023-07-27T09:43:42.4223852992 Request headers:
2023-07-27T09:43:42.4223928982 'X-IDENTITY-HEADER': 'REDACTED'
2023-07-27T09:43:42.4223986682 'User-Agent': 'azsdk-python-identity/1.13.0 Python/3.11.4 (Linux-5.15.0-1041-azure-x86_64-with-glibc2.36)'
2023-07-27T09:43:42.4224036782 No body was attached to the request
2023-07-27T09:43:47.2094416752 2023-07-27 09:43:47 - Response status: 200
2023-07-27T09:43:47.2094802912 Response headers:
2023-07-27T09:43:47.2094876162 'Content-Type': 'application/json; charset=utf-8'
2023-07-27T09:43:47.2094919932 'Date': 'Thu, 27 Jul 2023 09:43:46 GMT'
2023-07-27T09:43:47.2094962232 'Server': 'Kestrel'
2023-07-27T09:43:47.2095004042 'Transfer-Encoding': 'chunked'
2023-07-27T09:43:47.2095041842 'X-CORRELATION-ID': 'REDACTED'
2023-07-27T09:43:47.2097907702 2023-07-27 09:43:47 - DefaultAzureCredential acquired a token from ManagedIdentityCredential
2023-07-27T09:43:47.3883574762 2023-07-27 09:43:47 - message='OpenAI API response' path=https://blueopenai.openai.azure.com//openai/deployments/gpt-35-turbo-16k/chat/completions?api-version=2023-06-01-preview processing_ms=111.8376 request_id=7bee6697-a7a7-45a6-a782-a562a101f201 response_code=200
2023-07-27T09:44:48.27245 No logs since last 60 seconds
2023-07-27T09:45:48.68011 No logs since last 60 seconds

```

Alternatively, you can click open the **Logs** panel, as shown in the following screenshot, and use a [Kusto Query Language \(KQL\)](#) query to filter, project, and retrieve only the desired data.

Microsoft Azure (Preview) Search resources, services, and docs (G+)

Home > chatapp

chatapp | Logs

Container App

Search

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Application

Revisions

Containers

Scale and replicas

Settings

Authentication

Secrets

Ingress

Continuous deployment

Custom domains

Dapr

Identity

Service Connector (preview)

CORS (preview)

Locks

Monitoring

Alerts

Metrics

Logs

Log stream

Console

Advisor recommendations

Help

New Query 1*

BlueWorkspace Select scope

Run Time range: Last 24 hours Save Share New alert rule

Tables Queries Functions

Search

Filter Group by: Solution

Collapse all

Favorites

You can add favorites by clicking on the ☆ icon

LogManagement

- AzureDiagnostics
- AzureMetrics
- ContainerRegistryLoginEvents
- ContainerRegistryRepositoryEvents
- Usage

Custom Logs

- ContainerAppConsoleLogs_CL
- ContainerAppSystemLogs_CL

```

1 ContainerAppConsoleLogs_CL
2 | where ContainerAppName_s == 'chatapp'
3 | order by TimeGenerated desc

```





















Results Chart

TimeGenerated [UTC]	time_t [UTC]	_timestamp_d	ContainerName
> 7/27/2023, 9:50:30.254 AM	7/27/2023, 9:50:29.255 AM	1690451429.25546	chatapp
> 7/27/2023, 9:50:17.327 AM	7/27/2023, 9:50:16.346 AM	1690451416.34624	chatapp
> 7/27/2023, 9:50:16.343 AM	7/27/2023, 9:50:16.131 AM	1690451416.13101	chatapp
> 7/27/2023, 9:50:16.343 AM	7/27/2023, 9:50:16.131 AM	1690451416.131	chatapp
> 7/27/2023, 9:50:16.343 AM	7/27/2023, 9:50:16.130 AM	1690451416.13096	chatapp
> 7/27/2023, 9:43:48.366 AM	7/27/2023, 9:43:47.209 AM	1690451027.20944	chatapp
> 7/27/2023, 9:43:48.366 AM	7/27/2023, 9:43:47.209 AM	1690451027.20949	chatapp
> 7/27/2023, 9:43:48.366 AM	7/27/2023, 9:43:47.209 AM	1690451027.2095	chatapp
> 7/27/2023, 9:43:48.366 AM	7/27/2023, 9:43:47.209 AM	1690451027.20949	chatapp
> 7/27/2023, 9:43:48.366 AM	7/27/2023, 9:43:47.388 AM	1690451027.38836	chatapp
> 7/27/2023, 9:43:48.366 AM	7/27/2023, 9:43:47.209 AM	1690451027.2095	chatapp
> 7/27/2023, 9:43:48.366 AM	7/27/2023, 9:43:47.209 AM	1690451027.20948	chatapp
> 7/27/2023, 9:43:48.366 AM	7/27/2023, 9:43:47.209 AM	1690451027.2095	chatapp
> 7/27/2023, 9:43:48.366 AM	7/27/2023, 9:43:47.209 AM	1690451027.20979	chatapp
> 7/27/2023, 9:43:43.454 AM	7/27/2023, 9:43:42.422 AM	1690451022.42239	chatapp
> 7/27/2023, 9:43:43.454 AM	7/27/2023, 9:43:42.422 AM	1690451022.42233	chatapp
> 7/27/2023, 9:43:43.454 AM	7/27/2023, 9:43:42.421 AM	1690451022.42197	chatapp
> 7/27/2023, 9:43:43.454 AM	7/27/2023, 9:43:42.422 AM	1690451022.4224	chatapp
> 7/27/2023, 9:43:43.454 AM	7/27/2023, 9:43:42.422 AM	1690451022.42239	chatapp

0s 918ms Display time (UTC+00:00) Query details 1 - 20 of 23

Review deployed resources

You can use the Azure portal to list the deployed resources in the resource group, as shown in the following picture:

✓ Azure OpenAI				
<input type="checkbox"/>	 BlueOpenAI	Azure OpenAI	East US	...
✓ Container App				
<input type="checkbox"/>	 chatapp	Container App	East US	...
<input type="checkbox"/>	 docapp	Container App	East US	...
✓ Container Apps Environment				
<input type="checkbox"/>	 BlueEnvironment	Container Apps Environment	East US	...
✓ Container registry				
<input type="checkbox"/>	 BlueRegistry	Container registry	East US	...
✓ Log Analytics workspace				
<input type="checkbox"/>	 BlueWorkspace	Log Analytics workspace	East US	...
✓ Managed Identity				
<input type="checkbox"/>	 BlueRegistryIdentity	Managed Identity	East US	...
<input type="checkbox"/>	 BlueWorkloadIdentity	Managed Identity	East US	...
✓ Network Interface				
<input type="checkbox"/>	 BlueEnvironment	Container Apps Environment	East US	...
✓ Container registry				
<input type="checkbox"/>	 BlueRegistry	Container registry	East US	...
✓ Log Analytics workspace				
<input type="checkbox"/>	 BlueWorkspace	Log Analytics workspace	East US	...
✓ Managed Identity				
<input type="checkbox"/>	 BlueRegistryIdentity	Managed Identity	East US	...
<input type="checkbox"/>	 BlueWorkloadIdentity	Managed Identity	East US	...
✓ Network Interface				
<input type="checkbox"/>	 BlueOpenAIPrivateEndpoint.nic.7cde7fdf-5590-4db7-9e7b-91...	Network Interface	East US	...
<input type="checkbox"/>	 BlueRegistryPrivateEndpoint.nic.5767ff52-4ee8-484a-832d-d4...	Network Interface	East US	...
✓ Private DNS zone				
<input type="checkbox"/>	 privatelink.azurecr.io	Private DNS zone	Global	...
<input type="checkbox"/>	 privatelink.openai.azure.com	Private DNS zone	Global	...
✓ Private endpoint				
<input type="checkbox"/>	 BlueOpenAIPrivateEndpoint	Private endpoint	East US	...
<input type="checkbox"/>	 BlueRegistryPrivateEndpoint	Private endpoint	East US	...
✓ Virtual network				
<input type="checkbox"/>	 BlueVNet	Virtual network	East US	...

You can also use Azure CLI to list the deployed resources in the resource group:

```
1 | az resource list --resource-group <resource-group-name>
```

You can also use the following PowerShell cmdlet to list the deployed resources in the resource group:

```
1 | Get-AzResource -ResourceGroupName <resource-group-name>
```

Clean up resources

You can delete the resource group using the following Azure CLI command when you no longer need the resources you created. This will remove all the Azure resources.

```
1 | az group delete --name <resource-group-name>
```

Alternatively, you can use the following PowerShell cmdlet to delete the resource group and all the Azure resources.

```
1 | Remove-AzResourceGroup -Name <resource-group-name>
```

👍 3 Likes

6 Comments



[cicorias](#) Microsoft

Jul 27 2023 08:40 A

Thanks again for such relevant and timely guidance Paolo!!

👍 1 Like



[paolosalvatori](#) Microsoft

Jul 27 2023 09:01 A

Thanks [@cicorias](#), I thought to create an Azure Container Apps + Azure OpenAI sample after my articles on AKS + Azure OpenAI.

👍 1 Like



[merveguel](#) Copper Contributor

Aug 15 2023 02:32

Hi [@paolosalvatori](#), thanks for such detailed post, I was wondering if this article can be implemented using the free-tier version of the prerequisite tools?

Best, Merve

👍 0 Likes



[paolosalvatori](#) Microsoft

Aug 15 2023 03:17

[Skip to main navigation](#)

Hi @merveguel, all the pre-requisites are free of charge. I think you can use a free Azure account to test the entire architecture, but I'm not sure about Azure OpenAI Service. I should check, but I'm OOF now. If you don't have an Azure account already, the best way is to open one. Alternatively, check the pricing page for the various services and the Azure Free Account FAQ: <https://azure.microsoft.com/en-us/free/free-account-faq>

👍 1 Like



Heman_k Copper Contributor

Aug 29 2023 01:02

Hi @paolosalvatori, Thanks for sharing this very detailed article. Covers a lot of ground!!

On a related topic, can you please comment on the vector search capabilities in Azure Search? How does it compare it with vector search in other similar products?

Also, I am specifically interested in the support for HNSW algorithm in Azure Search. Does Azure Search support HNSW? I have seen sample code where there is a module called HnswVectorSearchAlgorithmConfiguration in the azure.search.documents.indexes.models package, but I am having problems importing this module, at least in Python 3.9. Any info or suggestions?

👍 0 Likes



paolosalvatori Microsoft

Aug 29 2023 11:59

Thanks @Heman_k

I'm not an Azure Cognitive Search subject-matter expert, but I can surely affirm that Azure Cognitive Search has strong vector search capabilities that allow you to search and retrieve similar vectors based on their similarity scores. While I do not have direct comparisons to other similar products or other vector databases, such as Chroma or FAISS, Azure Cognitive Search's vector search capabilities are designed to provide efficient and accurate results.

Regarding the HNSW algorithm, Azure Search does support it. The HNSW algorithm is a commonly used algorithm for approximate nearest neighbor search, and it is available in Azure Cognitive Search for vector search scenarios.

[Skip to main content](#)

Please check [Add vector search - Azure Cognitive Search | Microsoft Learn](#).
As for the specific issue you mentioned regarding importing the
HnswVectorSearchAlgorithmConfiguration module in Python 3.9, it's possible
that the module may not be available or accessible in that specific version. It is
recommended to check the documentation or reach out to Microsoft Azure
support for further assistance or alternative approaches.
Overall, Azure Search offers robust vector search capabilities, including support
for the HNSW algorithm, which can be valuable for creating efficient and
accurate search experiences.

P.S. If you found my article and sample interesting and helpful, please like the
article and star the GitHub project, thanks :)

 1 Like

You must be a registered user to add a comment. If you've already registered, sign in.
Otherwise, register and sign in.



[Comment](#)

Co-Authors



[paolosalvatori](#)

Version history

Last update: Jul 27 2023 07:40 AM
Updated by: [paolosalvatori](#)

Labels

App	101
Cloud Native Apps	70
Data & AI	110

< Previous

Next >

Share



What's new

Surface Pro 9

Surface Laptop 5

Surface Studio 2+

Surface Laptop Go 2

Surface Laptop Studio

Surface Duo 2

Microsoft 365

Windows 11 apps

Microsoft Store

Account profile

Download Center

Microsoft Store support

Returns

Order tracking

Virtual workshops and training

Microsoft Store Promise

Flexible Payments

Education

Microsoft in education

Devices for education

Microsoft Teams for Education

[Skip to footer navigation](#)

[Microsoft 365 Education](#)

[Education consultation appointment](#)

[Educator training and development](#)

[Deals for students and parents](#)

[Azure for students](#)

Business

[Microsoft Cloud](#)

[Microsoft Security](#)

[Dynamics 365](#)

[Microsoft 365](#)

[Microsoft Power Platform](#)

[Microsoft Teams](#)

[Microsoft Industry](#)

[Small Business](#)

Developer & IT

[Azure](#)

[Developer Center](#)

[Documentation](#)

[Microsoft Learn](#)

[Microsoft Tech Community](#)

[Azure Marketplace](#)

[AppSource](#)

[Visual Studio](#)

Company

[Careers](#)

[About Microsoft](#)

[Company news](#)

[Privacy at Microsoft](#)

[Investors](#)

[Diversity and inclusion](#)

[Accessibility](#)

[Skip to footer navigation](#)

